

Lecture 10-12: Language Models

USC VSoE CSCI 544: Applied Natural Language Processing


Jonathan May -- 梅約納

September 27-October 4, 2017

Quiz 1

- Which of the following CFG rules are in CNF (assume capital letters signify nonterminals, lowercase letters signify terminals)?
 - $S \rightarrow NP VP$
 - $NP \rightarrow NP PP$
 - $NP \rightarrow DT JJ NN$
 - $NP \rightarrow \text{the boy}$
 - $NP \rightarrow PP$
 - $JJ \rightarrow \text{red}$
 - $NP \rightarrow \text{the NP}$

Quiz 2



Step	Stack	Word List	Action	Relation Added
0	[root]	[book, me, the, morning, flight]	SHIFT	<div> <div>book</div> <div> <div>→</div> <div>me</div> </div> <div>dobj</div> </div>
1	[root, book]	[me, the, morning, flight]	SHIFT	
2	[root, book, me]	[the, morning, flight]	RArc-dobj	
3	[root, book]	[the, morning, flight]	SHIFT	
4	[root, book, the]	[morning, flight]	SHIFT	
5	[root, book, the, morning]	[flight]	SHIFT	
6	[root, book, the, morning, flight]	[]		

- What is the next step?

- shift
- **reduce**
- LArc
- RArc
- LArc-det

- RArc-det
- RArc-nmod
- LArc-nmod

Please turn your homework ...

- What word comes next?
 - in
 - over
 - into
 - the
 - refrigerator
- What are the probabilities of each of these?
- And why should we care?

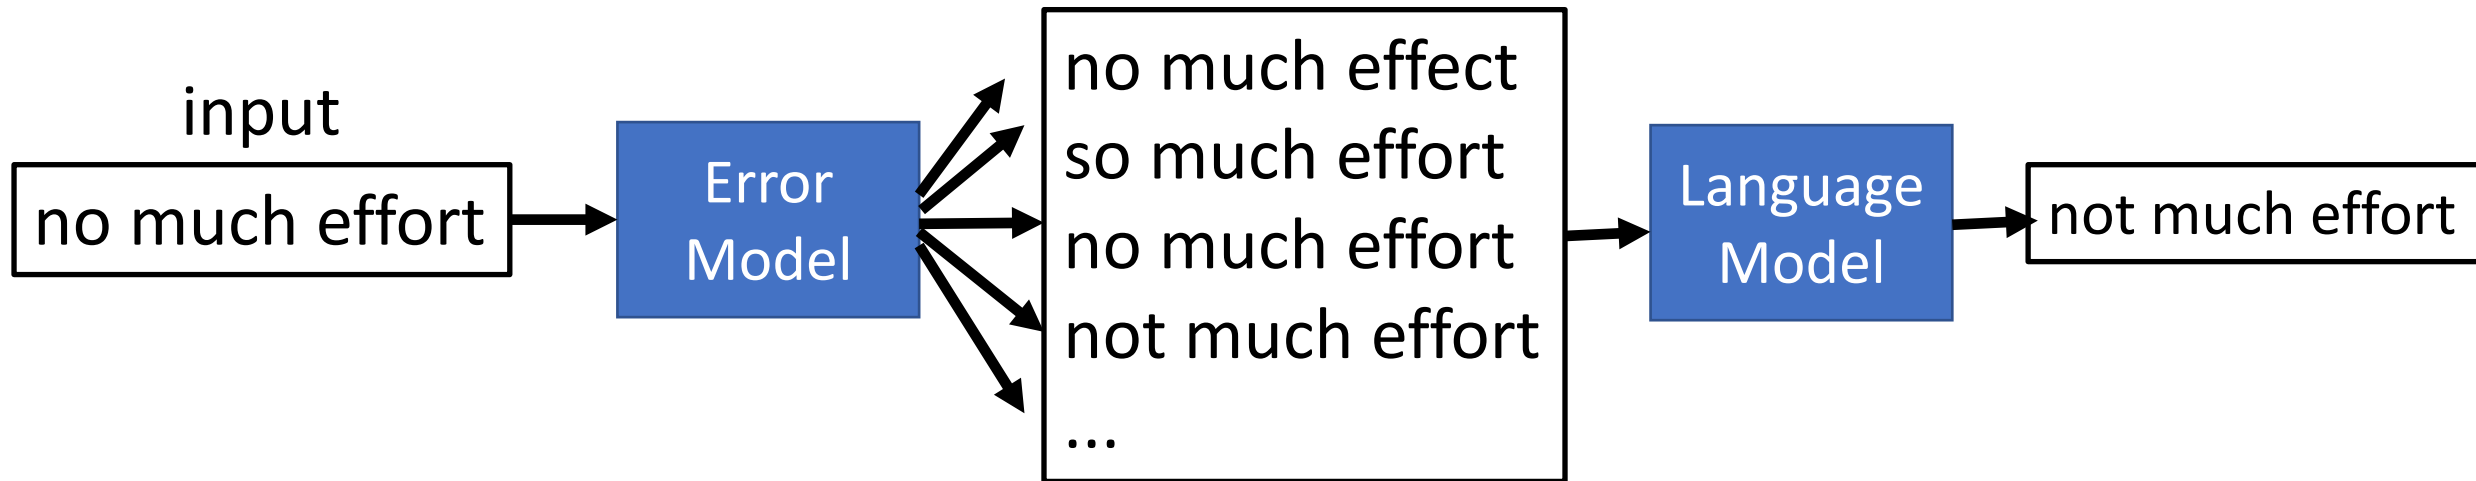
Probability of a sentence; $P(s)$

- How likely is it to occur?
- Colloquially, how likely is any speaker of language X to utter s?
 - $P(\text{the cat slept peacefully}) > P(\text{slept the peacefully cat})$
 - $P(\text{she studies morphosyntax}) > P(\text{she studies more faux syntax})$

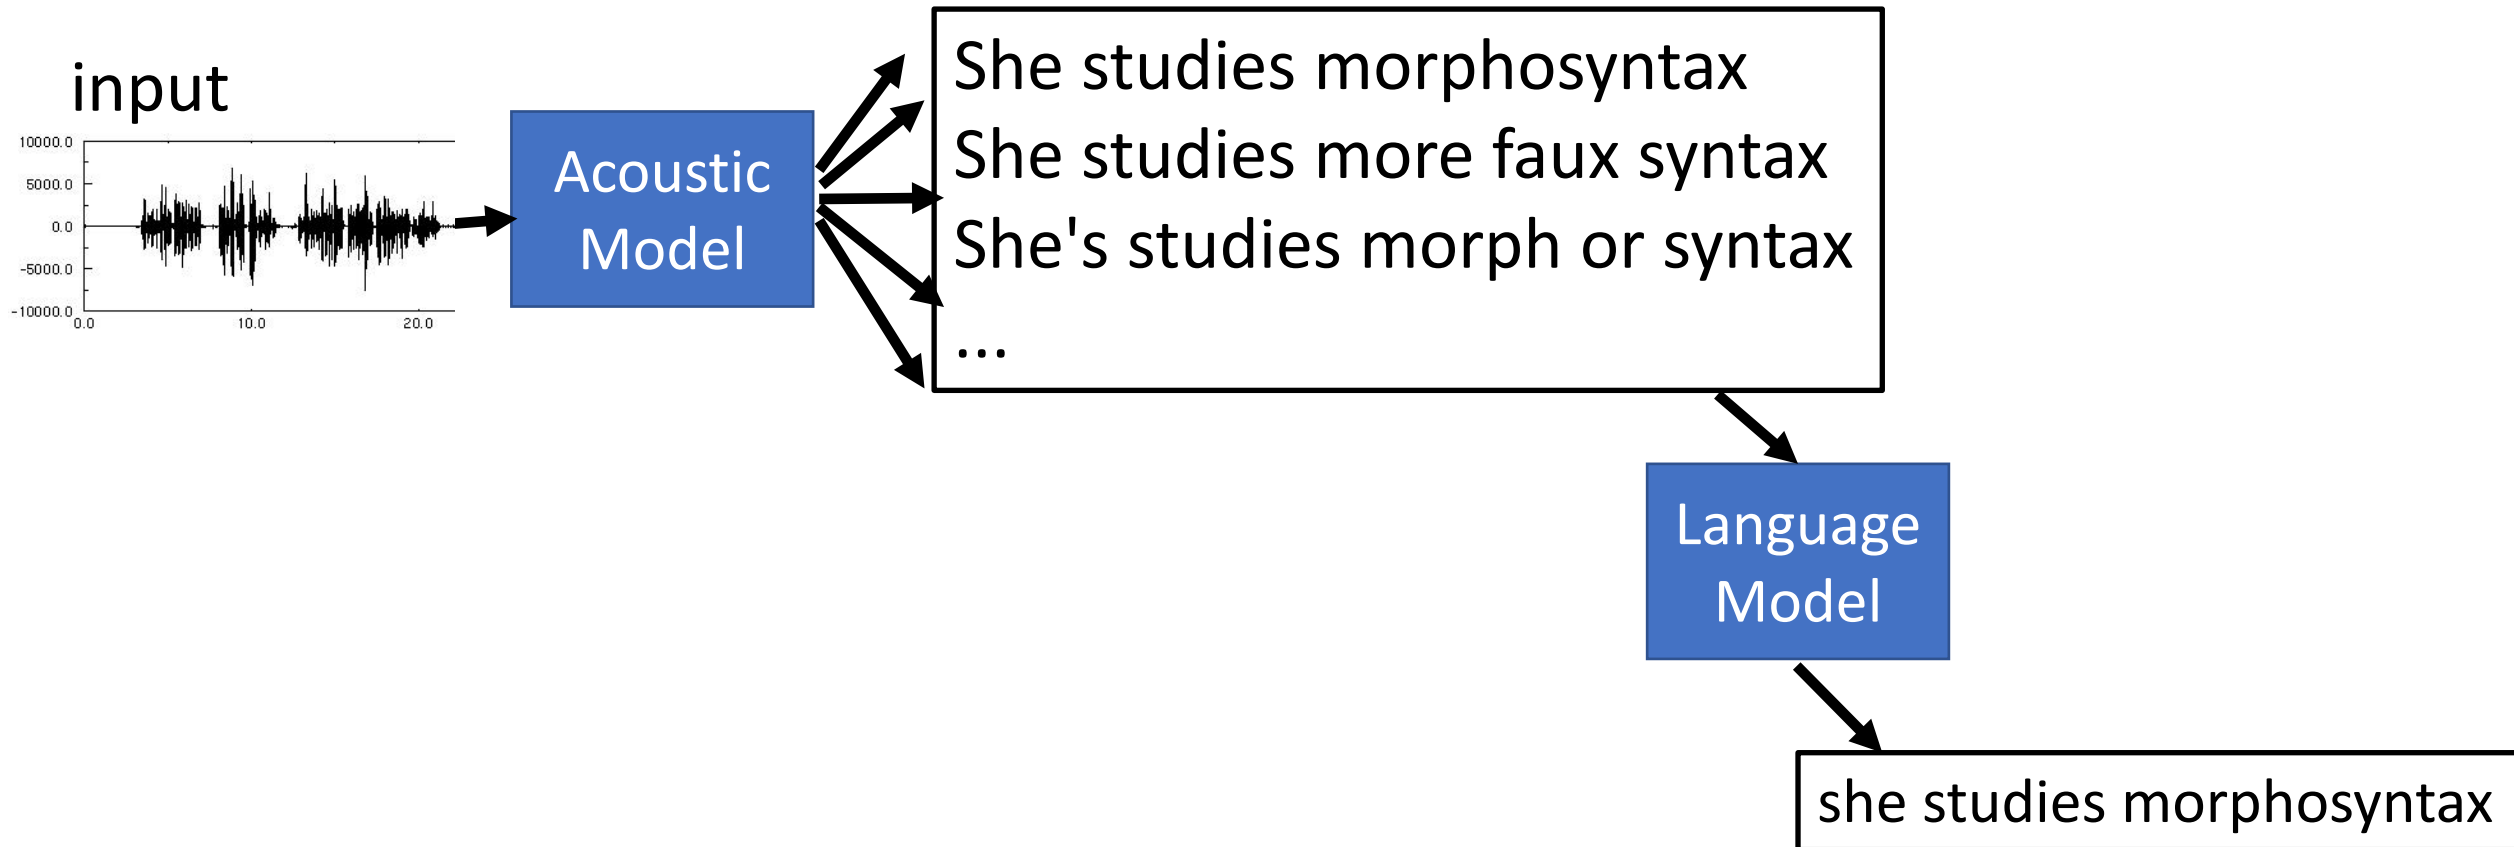
Language Models in NLP

- It's very difficult to know the true $P(s)$ for an arbitrary sequence of words
- But we can define a language model that will give us good approximations
- Language models (LMs) are very useful whenever we are generating output
 - Machine Translation
 - Spelling Correction
 - Summarization
 - Speech Recognition
- There are some very good, easy-to-use toolkits for building and using LMs
 - SRILM: around since the 90s. not advised with >300m tokens
 - KenLM: preferred choice; great scalability in memory and time with even billions of tokens
 - NLTK has some LM training/using support (ok for prototyping)

Use of Language Models: Spelling Correction



Use of Language Models: Speech Recognition



Use of Language Models: Machine Translation

input

ella se va a casa

Translation
Model

She is going home
She is going house
She goes to home
To home she is going
...

Language Model

she is going home

```
graph LR; Input[ella se va a casa] --> TM[Translation Model]; TM --> Candidates["She is going home<br/>She is going house<br/>She goes to home<br/>To home she is going<br/>..."]; Candidates --> LM[Language Model]; LM --> Output[she is going home]
```

The diagram illustrates the Machine Translation process. It starts with an input box containing the Spanish phrase "ella se va a casa". An arrow points from this box to a blue box labeled "Translation Model". From the "Translation Model" box, four arrows point to a larger box containing four candidate translations: "She is going home", "She is going house", "She goes to home", and "To home she is going", followed by an ellipsis "...". An arrow then points from this candidate box to another blue box labeled "Language Model". Finally, an arrow points from the "Language Model" box to an output box containing the selected translation: "she is going home".

LMs for Prediction

- LMs can be used to predict what a human will do next, rather than correct a possibly faulty model
- Example: predictive text correction/completion on your phone
 - Keyboard is tiny, so it's easy to touch a spot slightly off from the letter you intend
 - Correct these errors as you go and also provide possible completions

i n e f f i c i e n t

- In this case, LM may be defined over sequences of *characters* instead of (or in addition to) sequences of words

Noisy Channel Model



But How To Estimate These Probabilities?

- We want to know the probability of word sequence $\mathbf{w} = w_1, w_2, \dots, w_n$ occurring in English
- Assume we have some training data: large corpus of general English text
- We use this data to estimate the probability of \mathbf{w} (even if we never see it in the corpus)

Bit of Notation: Random Variables

- Random Variable = variable that represents all the possible events in some partition of Ω
- So if X is a coin flip I can say $P(X=\text{heads})$ to mean probability the flip comes up heads or just $P(X)$ to mean the probability table for the events (heads, tails)
- We can treat Random Variables like events
 - Flip coin A. Independently, flip coin B
 - $P(A=\text{heads} \mid B=\text{tails}) = P(A=\text{heads})$; $A=\text{heads}$ and $B=\text{tails}$ are independent events
 - for all x in {heads, tails}, for all y in {heads, tails}, $P(A=x \mid B=y) = P(A=x)$; A and B are independent random variables
- We've been using Random Variables all along, actually, but have been a bit sloppy
- (I probably should have discussed these back in Lecture 4)

Probability of a word sequence

- $P(\mathbf{w}) = P(w_1, w_2, w_3, w_4, \dots, w_n)$
- e.g. $P(\mathbf{w} = \text{the cat slept quietly}) = P(w_1=\text{the}, w_2=\text{cat}, w_3=\text{slept}, w_4=\text{quietly})$
- We'll often abuse notation when talking about specific events and context is clear, e.g. $P(\text{the cat slept quietly})$

Maximum Likelihood Estimation?

- Recall Maximum Likelihood Estimations (MLE) for our HMM POS tagger
 - AKA "Count and divide"
- So get a corpus of N sentences
 - $P_{MLE}(\mathbf{w} = \text{the cat slept quietly}) = C(\text{the cat slept quietly})/N$
- But consider these sentences:
 - the long-winded peripatetic beast munched contentedly on mushrooms
 - parsimonius caught the of about for syntax
- Neither is in a corpus (I just made them up), so $P_{MLE}=0$ for both
 - But one is meaningful and grammatical and the other isn't!

Sparse Data and MLE

- If something doesn't occur, MLE thinks it can't occur
- No matter how much data you get, you won't have enough observations to model all events well with MLE
- We need to make some assumptions so that we can provide a reasonable probability for grammatical sentences, even if we haven't seen them

Independence (Markov) Assumption

- Recall, $P(w_1, w_2, \dots, w_n) = P(w_n | w_1, w_2, \dots, w_{n-1})P(w_{n-1} | w_1, w_2, \dots, w_{n-2})\dots$
 - $\prod_{i=1}^n P(w_i | w_1, \dots, w_{i-1})$
- Still too sparse (nothing changed; same information)
 - if we want $P(\text{I spent three years before the mast})$
 - we still need $P(\text{mast} | \text{I spent three years before the})$
- Note: could use chain rule any number of ways
 - $P(w_4 = \text{years} | w_1 = \text{I}, w_2 = \text{spent}, w_3 = \text{three}, w_5 = \text{before}, w_6 = \text{the}, w_7 = \text{mast})^* \dots$
- Remember definition of independence; A and B are independent if $P(A) = P(A|B)$

Independence (Markov) Assumption

- Make an n-gram independence assumption: probability of a word only depends on a fixed number of previous words (history)
 - **trigram model**: $P(w_i | w_1, \dots, w_{i-1}) \approx P(w_i | w_{i-2}, w_{i-1})$
 - **bigram model**: $P(w_i | w_1, \dots, w_{i-1}) \approx P(w_i | w_{i-1})$
 - **unigram model**: $P(w_i | w_1, \dots, w_{i-1}) \approx P(w_i)$
- I.e. a trigram model says
 - $P(\text{mast} \mid \text{I spent three years before the}) \approx P(\text{mast} \mid \text{before the})$
- It also assumes all these are equal:
 - $P(\text{mast} \mid \text{I spent three years before the})$
 - $P(\text{mast} \mid \text{I went home before the})$
 - $P(\text{mast} \mid \text{I saw the sail before the})$because all are estimated as $P(\text{mast} \mid \text{before the})$
- Not always a good assumption! But it does reduce the sparse data problem

Estimating Trigram Conditional Probabilities

- $P_{MLE}(\text{mast} \mid \text{before the}) = \text{Count}(\text{before the mast}) / \text{Count}(\text{before the})$

- In general, for any trigram, we have

- $P_{MLE}(w_i \mid w_{i-2}, w_{i-1}) = \frac{\text{Count}(w_{i-2}, w_{i-1}, w_i)}{\text{Count}(w_{i-2}, w_{i-1})}$

- To be clear, the MLE uses all data, not just data for the particular random variables we're estimating.

- $\text{Count}(\text{"the sequence before the mast"})$, not $\text{Count}(\text{"the sequence where } w_5=\text{before}, w_6=\text{the}, w_7=\text{mast}"})$

Example from *Moby Dick* corpus

- $C(\text{before, the}) = 25$; $C(\text{before, the, mast}) = 4$
- $C(\text{before, the, mast}) / C(\text{before, the}) = 0.16$
- mast is the most common word to come after "before the" (wind is second most common)
- $P_{\text{MLE}}(\text{mast}) = 56/110927 = .0005$ and $P_{\text{MLE}}(\text{mast}|\text{the}) = .003$
- Seeing "before the" vastly increases the probability of seeing "mast" next

Trigram model summary

- To estimate $P(\mathbf{w})$, use chain rule and make an independence assumption
 - $P(w_1, \dots, w_n) = \prod_{i=1}^n P(w_i | w_1, \dots, w_{i-1})$
 - $\approx P(w_1)P(w_2|w_1) \prod_{i=3}^n P(w_i|w_{i-2}, w_{i-1})$
- Then estimate each trigram prob from data (here, using MLE)
 - $P_{MLE}(w_i | w_{i-2}, w_{i-1}) = \frac{\text{Count}(w_{i-2}, w_{i-1}, w_i)}{\text{Count}(w_{i-2}, w_{i-1})}$

Midterm Info

Some of you will
NOT BE IN SAL 101
YOU WILL BE
IN MHP (Mudd) 101



Who is Where?

- If your last* name is from Aggarwal to Patel, you are in SAL 101
- If it is from Pathapi to Zhu, you are in MHP 101
- * "last" is a very Eurocentric term. I have an alphabetization based on what USC has your surname listed as.
- To avoid confusion, please check the course website, which has a very prominent link to a roster with firstname, last name, partially masked email address, and assignment
- If you are still unsure of where to go, contact us on piazza ASAP

Midterm Logistics

- Length: 1 hour, 40 minutes (you will need time to set up and hand in exam)
- Date: Friday, October 6, 8:00 AM (Please arrive promptly!)
- Please Bring:
 - Pencils/pens/erasers as needed
 - one 8.5x11 inch (or A4) sheet of paper with notes on both sides (optional)
 - NO OTHER NOTES
 - NO ELECTRONIC RESOURCES
 - NO BOOKS
- We will provide extra paper for scratch work
- Sit only at seats with exams on them. Fill up all available space.

What's On The Exam?

- Fair game
 - Anything on the slides
 - Anything in the required reading
 - Anything in the homeworks
- But
 - We're not trying to trick you
 - We're not trying to make this impossible
 - If you understand the lectures well, you should be ok

What's On The Exam

- Major Topics
 - Levels of Linguistic Knowledge (L1)
 - Corpora, Regex, Basic text processing (L2)
 - Morphology, Finite State Automata and Transducers (L3)
 - Probability Theory (L4)
 - Naive Bayes, Features, Perceptron, Logistic Regression (L5)
 - POS Tagging and HMM tagger, Viterbi Decoding (L6)
 - Constituency Syntax Trees, Context-Free Grammars, CKY, CNF, Smoothing, Interpolating, Beam Decoding (L7-8)
 - Dependency Syntax Trees, Arc-Standard and Arc-Eager Dependency Parsing (L9-10)
 - Ngram Language Models, Smoothing, Backoff, Interpolation, alternate language models (L10-11) (Probably no neural; depends on how far we get)
- It won't all be on there because there isn't enough time
 - But there is plenty of room on the final

Practical details (I)

- Trigram model assumes two-word history
- But consider these sentences:

w_1	w_2	w_3	w_4
he	saw	the	yellow
feeds	the	cats	daily

- What's wrong?
 - a sentence shouldn't end with 'yellow'
 - a sentence shouldn't begin with 'feeds'
- Does the model capture these problems?

Beginning / end of sequence

- To capture behavior at beginning/end of sequences, we can augment the input:

w_{-1}	w_0	w_1	w_2	w_3	w_4	w_5
<s>	<s>	he	saw	the	yellow	</s>
<s>	<s>	feeds	the	cats	daily	</s>

- That is, assume $w_{-1}=w_0=<s>$ and $w_{n+1}=</s>$ so:
 - $P(\mathbf{w}) = \prod_{i=1}^{n+1} P(w_i | w_{i-2}, w_{i-1})$
- Now $P(</s> | \text{the, yellow})$ is low, indicating this is not a good sentence
- $P(\text{feeds} | <s>, <s>)$ should also be low

Beginning/end of sequence

- Alternatively, we could model all sentences as one (very long) sequence, including punctuation
 - two cats live in sam 's barn . sam feeds the cats daily . yesterday , he saw the yellow cat catch a mouse . [...]
- Now, trigram probabilities like $P(. \mid \text{cats daily})$ and $P(, \mid . \text{ yesterday})$ tell us about behavior at sentence edges
- Here, all tokens are lowercased. What are the pros/cons of not doing that?

Practical details (II)

- Word probabilities are typically very small.
- Multiplying lots of small probabilities quickly gets so tiny we can't represent the numbers accurately, even with double precision floating point.
- So in practice, we typically use log probabilities (usually base-e)
 - Since probabilities range from 0 to 1, log probs range from $-\infty$ to 0
 - Instead of multiplying probabilities, we add log probs
 - Often, negative log probs are used instead; these are often called "costs"; lower cost = higher prob
- Recall: we saw this with bigram HMM for POS tagging

Interim Summary: N-gram probabilities

- "Probability of a sentence": how likely is it to occur in natural language?
- We can never know the true probability, but we may be able to estimate it from corpus data.
- N-gram models are one way to do this:
 - To alleviate sparse data, assume word probs depend only on short history
 - Tradeoff: longer histories may capture more, but are also sparser
 - So far, we estimated N-gram probabilities using MLE

Interim Summary: Language Models

- Language Models tell us $P(\mathbf{w}) = P(w_1, \dots, w_n)$: How likely is this sequence of words to occur?
 - Roughly: Is this sequence of words a "good" one in my language?
- LMs are used as a component in applications such as speech recognition, machine translation, and predictive text completion
- To reduce sparse data, N-gram LMs assume words depend only on a fixed-length history, even though we know this isn't true
- Next:
 - How to evaluate a language model
 - Weaknesses of MLE and how to address them (more sparsity)

Quiz 3

- Which of These Statements are True?
 - Naive Bayes is a probabilistic model
 - Perceptron is a generative model
 - Logistic regression has a closed-form solution
 - Logistic regression is a discriminative model

Two Types of Evaluation in NLP

- **Extrinsic:** measure performance on a downstream application
 - For LM, plug it into a machine translation/ASR/etc system
 - The most reliable and useful evaluation: We don't use LMs absent other technology
 - But can be time-consuming
 - And of course we still need an evaluation measure for the downstream system
- **Intrinsic:** design a measure that is inherent to the current task
 - much quicker/easier during development cycle
 - not always easy to figure out what the right measure is. Ideally, it's one that correlates with extrinsic measures
 - Extra-hard for LMs

Intrinsically Evaluating a Language Model

- For parsing, tagging, sentiment, etc. it was fairly clear how to evaluate: Hold a set of labeled data out and see how often your model gets it right
- For LM, it's not quite so clear
 - Given a corpus of sentences and non-sentences, see how often the LM thinks you have a sentence?
 - Not a very realistic evaluation of how an LM is used
 - Often we are deciding between not-that-grammatical outputs
- Ideally we want a regression evaluation
 - Given a sentence, how close is the model probability to the true probability
 - But we don't know the true probability of a sentence!

Idea: Model should give high probability to an unseen corpus

- Assume that you have a proper probability model, i.e. for all sentences S in the language L , $\sum_{S \in L} P(S) = 1$
- Then take a held-out test corpus T consisting of sentences in the language you care about
- $\prod_{t \in T} P(t)$ should be as high as possible; model should think each sentence is a good one
- Let's be explicit about evaluating each word in each sentence
 - $\prod_{t \in T} \prod_{w \in t} P(w)$
- Collapse all these words into one big 'sentence' N :
 - $\prod_{w \in N} P(w)$

Resolving Some Problems

- $\prod_{w \in N} P(w)$ is going to result in underflow. Ok, let's use logs again!
- Also we tend to like positive sums.
 - $-\sum_{w \in N} \log_2 (P(w))$
- This can be tough to compare against corpora of different length (or sentences of different length), so normalize by the number of words:
 - $\frac{-\sum_{w \in N} \log(P(w))}{|N|}$ is called the cross-entropy of the data according to the model
- When comparing models, differences between these numbers tend to be pretty small, so we exponentiate
 - $2^{\frac{-\sum_{w \in N} \log(P(w))}{|N|}}$ is called the perplexity of the data
- Think of this as "how surprised is the model?"

Example

- Three word sentence with probabilities $\frac{1}{4}$, $\frac{1}{2}$, $\frac{1}{4}$
 - $\frac{1}{4} * \frac{1}{2} * \frac{1}{4} = .03125$
 - cross-entropy: $-(\log(1/4) + \log(1/2) + \log(1/4))/3 = 5/3$; $2^{5/3} \approx 3.17$
- Six word sentence with probabilities $\frac{1}{4}$, $\frac{1}{2}$, $\frac{1}{4}$, $\frac{1}{4}$, $\frac{1}{2}$, $\frac{1}{4}$
 - $\frac{1}{4} * \frac{1}{2} * \frac{1}{4} * \frac{1}{4} * \frac{1}{2} * \frac{1}{4} = .00097$
 - cross-entropy: $-(\log(1/4) + \log(1/2) + \log(1/4) + \log(1/4) + \log(1/2) + \log(1/4))/6 = 10/6$; $2^{10/6} \approx 3.17$
- If you overfit your training corpus so that $P(\text{train}) = 1$, then Perplexity on train is 0
- But Perplexity on test (which doesn't overlap with train) will be infinite

Intrinsic Evaluation Big Picture

- Lower Perplexity is better
- Roughly = number of bits needed to communicate information about a word
 - The terms 'cross-entropy' and 'perplexity' come out of information theory; it's in the reading if you're interested but we won't dwell on it
- In principle you could compare on different test sets
- In practice, domains shift. To know which of two LMs is better, train on common training sets, test on common test sets

Sparse data, again

- Suppose we build a trigram model from Moby Dick and evaluate the sentence "I spent three years before the mast"
- "I spent three" never occurs in training, so $P_{MLE}(\text{three} | \text{I spent}) = 0$
- so cross-entropy is infinite
- This is basically right; our model says "I spent three" should never occur so when it does our model is infinitely surprised!
- Even with a unigram model we run into words we never saw, so we need better ways to estimate probabilities from sparse data

Add-1 (Laplace) and Add- α (Lidstone)

Smoothing Again

- Pretend we saw everything 1 (α) more times than we did before
- $P_{+1}(w_i | w_{i-2}, w_{i-1}) = (C(w_{i-2}, w_{i-1}, w_i) + 1) / (C(w_{i-2}, w_{i-1}) + |V|)$ where $|V|$ is the size of the vocabulary
- $P_{+\alpha}(w_i | w_{i-2}, w_{i-1}) = (C(w_{i-2}, w_{i-1}, w_i) + \alpha) / (C(w_{i-2}, w_{i-1}) + \alpha |V|)$

Dealing with unknown vocabulary

- Can we also add a new 'OOV' token as was done in HMM emission table?
- It gets kind of complicated...
- $P_{+\alpha}(w_i | w_{i-2}, w_{i-1}) = (C(w_{i-2}, w_{i-1}, w_i) + \alpha) / (C(w_{i-2}, w_{i-1}) + \alpha |V| + 1)$
- $P_{+\alpha}(w_i = \text{OOV} | w_{i-2}, w_{i-1}) = \alpha / (\alpha |V| + 1)$
- But then we also have to deal with, e.g., $P_{+\alpha}(w_i | w_{i-2}, w_{i-1} = \text{OOV})$
- Better solution: replace low-count words in corpus with "OOV"
- Intuition: 1-count is basically the same as 0-count

Remaining Problem

- In a training corpus, suppose we see *Scottish beer* but neither of
 - *Scottish beer drinkers*
 - *Scottish beer eaters*
- If we build a smoothed trigram model (with any kind of smoothing), which example has higher probability?
 - Both the same! Unknown events are treated equally by smoothing!

Remaining Problem

- Previous smoothing methods assign equal probability to unseen events
- Better: use information from lower-order N-grams (shorter histories)
 - beer drinkers
 - beer eaters
- Two ways: backoff and interpolation

Backoff

- Idea: Trust the highest order language model that contains your N-gram
- $P_{BO}(z|x\ y) =$
 $(1 - \alpha_{xy})P(z|x\ y)$ if $\text{count}(x\ y) > 0$
 $\alpha_{xy} P_{BO}(z | y)$ else
- where α_{xy} is an interpolation parameter

Simple Interpolation

- Idea: Trust different amounts of context differently
- $P_{SI}(z|x y) =$
 $\lambda_3 P(z|x y) +$
 $\lambda_2 P(z|y) +$
 $\lambda_1 P(z) +$
 λ_0
- where $\lambda_0 + \lambda_1 + \lambda_2 + \lambda_3 = 1$, all ≥ 0
- We did something similar in lexicalized constituency parsing


Better Interpolation

- Idea: As in backoff, particular contexts matter
- $P_{SI}(z|x y) =$
 $\lambda_{xy} P(z|x y) +$
 $\lambda_{y/xy} P(z|y) +$
 $\lambda_{1/xy} P(z) +$
 $\lambda_{0/xy}$
- where, for each xy , $\lambda_{0/xy} + \lambda_{1/xy} + \lambda_{y/xy} + \lambda_{xy} = 1$, all ≥ 0
- Best not to actually have a different set for each unique context; can group by context count

State-of-the-art Smoothing

- There is lots and lots of work done on smoothing and lots of variants
- See Chen and Goodman (optional reading); it's actually quite comprehensive, though mathy
- Best today is Modified Kneser-Ney
 - replace MLE with estimates based on count of unique histories
 - 4 interpolation lambdas based on ngram counts
- For very large data, Google's Stupid Backoff
 - Really fast to calculate; good for very large data
 - Doesn't give proper perplexities!
 - Works well in practice
- These are available in SRILM (K-N) and KenLM (both)

Ngram LM as FSA



Quiz 4

stack	symbols
[root] a b c e f g	j k l o p

- The above table represents a potential configuration during dependency parsing. If we are doing **arc-eager** parsing, which symbols may be combined in a Left-Arc from this configuration?
 - f and g
 - e and f
 - a and o
 - g and j

Quiz 5

stack	symbols
[root] a b c e f g	j k l o p

- The above table represents a potential configuration during dependency parsing. If we are doing **arc-standard** parsing, which symbols may be combined in a Left-Arc from this configuration?
 - f and g
 - e and f
 - a and o
 - g and j

Other Approaches To Language Modeling

- Maybe we don't always care about the most immediate words
 - "Show Sally a good ____"
 - "Show Dave a good ____"
- $P(\text{time} \mid \text{a good})$ isn't so great
- $P(\text{time} \mid \text{Sally a good})$ isn't really either
- $P(\text{time} \mid \text{Show Sally a good})$ isn't helpful for Dave (and backoff/smoothing doesn't really help)
 - But how much to skip?
 - We could interpolate different skip models; doesn't help that much though

Other Approaches To Language Modeling

- Class-based smoothing
 - Train =party on Tuesday ...
 - Test =party on Monday ...
.....celebration on Tuesday...
- Maybe we could predict classes first, and then words...
 - Model $P(w_i = \text{Monday} \mid w_{i-2} = \text{party}, w_{i-1} = \text{on})$
 - $P(c_i = \text{DAY} \mid w_{i-2} = \text{party}, w_{i-1} = \text{on}) * P(w_i = \text{Monday} \mid w_{i-2} = \text{party}, w_{i-1} = \text{on}, c_i = \text{DAY})$
 - Or even drop the words themselves from the conditional
 - $P(c_i = \text{DAY} \mid w_{i-2} = \text{EVENT}, w_{i-1} = \text{PREP}) * P(w_i = \text{Monday} \mid c_i = \text{DAY})$
- These generally perform worse than trigram on their own but can help when interpolated

Other Approaches To Language Modeling

- Language should be syntactically well formed!
 - Can use law of total probability in reverse:
 - $P(s) = \sum_{t \in T} P(s, t)$ where t is a syntax tree for s
 - Or, instead of going in n-gram order, why not follow dependency links
 - $P(\text{I like your shoes}) = P(\text{like} | \text{root})P(\text{I} | \text{like})P(\text{shoes} | \text{like})P(\text{your} | \text{shoes})$
- Difficulties: Requires syntactic analysis which means less data available
- In practice, these methods, when interpolated with 3grams, helped a bit
- once we got beyond 1b words of data, not helpful

Adding Hidden Information

- Maybe there are different types of sentences!
 - Introductory
 - Asides
 - Technical
- And each one behaves differently!
- If you knew the type of the sentences you could evaluate it with a separate LM estimated off of just that type of sentence

Feature-Based LM

- Let's return to perceptron/maxent
- Previously we predicted sentiment, word sense, author given an input text
- Up to now we've been talking about predicting x_i given x_{i-1} , x_{i-2}
- Can we model this as a discriminative feature-based model?

	$\phi(x)$	$w(x_i=\text{bit})$	$w(x_i=\text{bought})$
bias	1	.95	-3.2
$x_{i-2}=\text{apple}$	0	.436	33.6
$x_{i-2}=\text{dog}$	0	-34	...
$x_{i-2}=\text{cat}$	0
$x_{i-2}=\text{the}$	1	3.4	3.6
...
$x_{i-1}=\text{apple}$	0
$x_{i-1}=\text{dog}$	1	14.4	-5.6
$x_{i-1}=\text{cat}$	0	6.3	-7.8
$x_{i-1}=\text{the}$	0	-5	-17

$P(\text{bit} \mid \text{the cat})$ vs $P(\text{bought} \mid \text{the cat})$


- Kind of like a trigram model!
- Note, separate weight for each output type again (too many to evaluate each)
- Bias term: a priori preference for that word (unigram prob)

	$\phi(x)$	$w(x_i=\text{bit})$	$w(x_i=\text{bought})$
bias	1	.95	-3.2
$x_{i-2}=\text{apple}$	0	.436	33.6
$x_{i-2}=\text{dog}$	0	-34	...
$x_{i-2}=\text{cat}$	0
$x_{i-2}=\text{the}$	1	3.4	3.6
...
$x_{i-1}=\text{apple}$	0
$x_{i-1}=\text{dog}$	1	14.4	-5.6
$x_{i-1}=\text{cat}$	0	6.3	-7.8
$x_{i-1}=\text{the}$	0	-5	-17
...	
$x_{i-2}=\text{NOUN}$	0
$x_{i-2}=\text{DET}$	1
$x_{i-1}=\text{NOUN}$	1
$x_{i-1}=\text{DET}$	0
$x_{i-2}=\text{apple} \wedge x_{i-1}=\text{the}$	1
	


$P(\text{bit} \mid \text{the cat})$ vs $P(\text{bought} \mid \text{the cat})$

- As before, we can add arbitrary features
 - POS tags
 - word length
 - initial letter
- What about conjuncts of features (true trigram)?
 - Yes, but this will get very lengthy
 - Shouldn't the individual units be enough?

Linear models can only separate linearly



Good Case



Need to add a dimension
with the conjunction of these features!

Bad Case

This May Seem Like A Weird Aside!

Don't worry, here's the road map:

N-gram language models ->

Considering other properties ->

Arbitrary feature language models ->

Handling conjuncts of features ->

Stacks of nonlinear perceptron = neural network language models

Solving the XOR problem

- Let our features be two dimensional input and the class label one-dimensional output
- Can we find W (2x1) and b (2x1) such that the function $y = Wx + b$ solves the problem for this data?
- No
- But we can use a nonlinear function and map this data into a new, separable feature space!

x1	x2	y
1	1	1
-1	1	-1
-1	-1	1
1	-1	-1

Mapping into a new space

h1 mapper

x1	1
x2	1
b	-1

x-space data

	x1	x2	y
a	1	1	1
b	-1	1	-1
c	-1	-1	1
d	1	-1	-1

h2 mapper

x1	-1
x2	-1
b	-1

$$\text{step}(x1 \cdot h_{x1} + x2 \cdot h_{x2} + b)$$


$$a: \text{step}(1+1-1) = 1$$

$$b: \text{step}(-1+1-1) = -1$$

$$c: \text{step}(-1-1-1) = -1$$

$$d: \text{step}(1-1-1) = -1$$

step function: 1 if > 0, else -1




$$a: \text{step}(-1-1-1) = -1$$

$$b: \text{step}(1-1-1) = -1$$


$$c: \text{step}(1+1-1) = 1$$

$$d: \text{step}(-1+1-1) = -1$$


x space



h space



New ('hidden') space to output space



	h1	h2	y
a	1	-1	1
b	-1	-1	-1
c	-1	1	1
d	-1	-1	-1

o mapper

h1	1
h2	1
b	1


$$\text{step}(h1 \cdot o_{h1} + h2 \cdot o_{h2} + b)$$

$$a: \text{step}(1 - 1 + 1) = 1$$

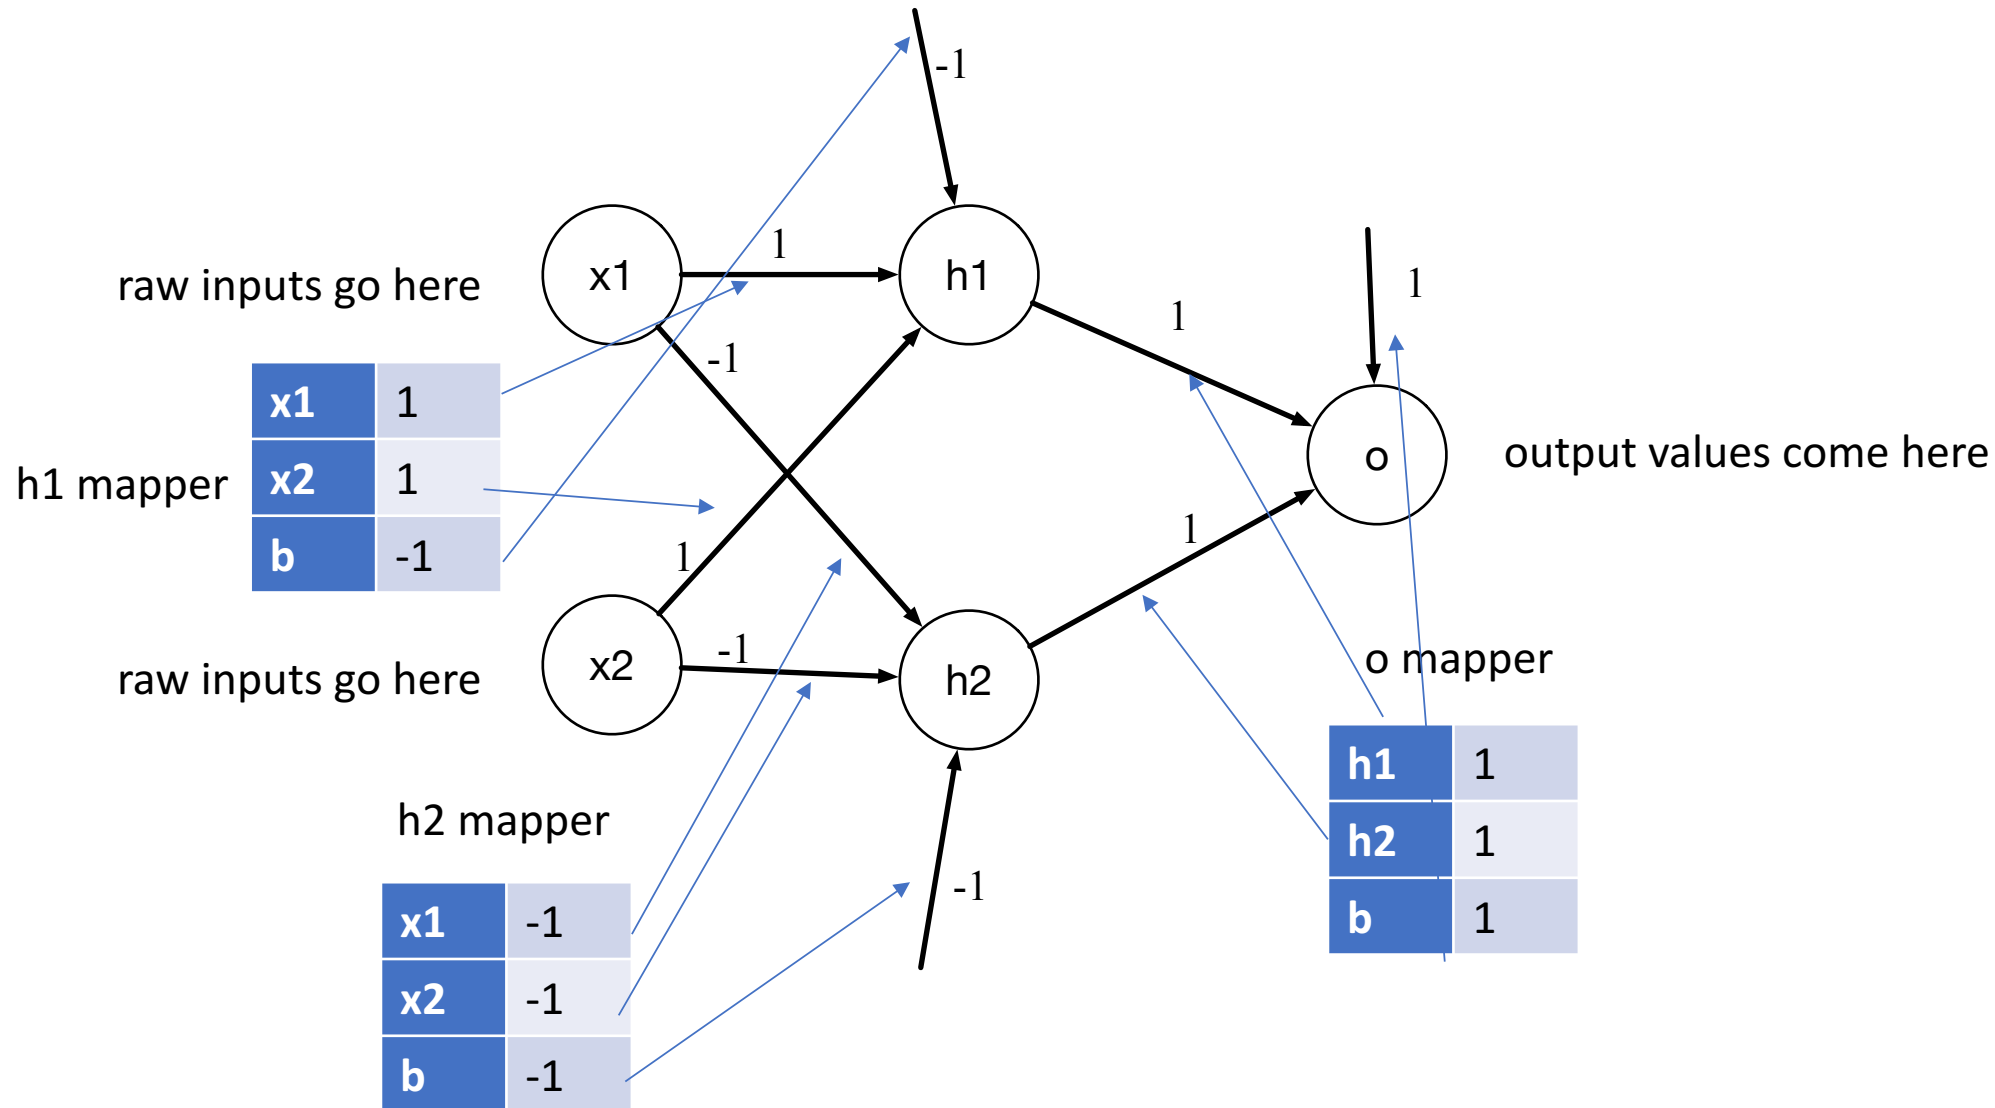
$$b: \text{step}(-1 - 1 + 1) = -1$$

$$c: \text{step}(-1 + 1 + 1) = 1$$

$$d: \text{step}(-1 - 1 + 1) = -1$$



Another way of looking at these matrices



Quiz 6

stack	symbols
[root] a b c e f g	j k l o p

- What is the resulting configuration after an arc-standard Left-Arc?

stack	symbols
[root] a b c e f	g j k l o p

stack	symbols
[root] a b c e f	j k l o p

stack	symbols
[root] a b c e g	j k l o p

stack	symbols
[root] a b c f g	j l o p

stack	symbols
[root] a b c e f g j	k l o p

Quiz 7

stack	symbols
[root] a b c e f g	j k l o p

- What is the resulting configuration after an arc-eager Right-Arc?

stack	symbols
[root] a b c e f	g j k l o p

stack	symbols
[root] a b c e f	j k l o p

stack	symbols
[root] a b c e g	j k l o p

stack	symbols
[root] a b c f g	j l o p

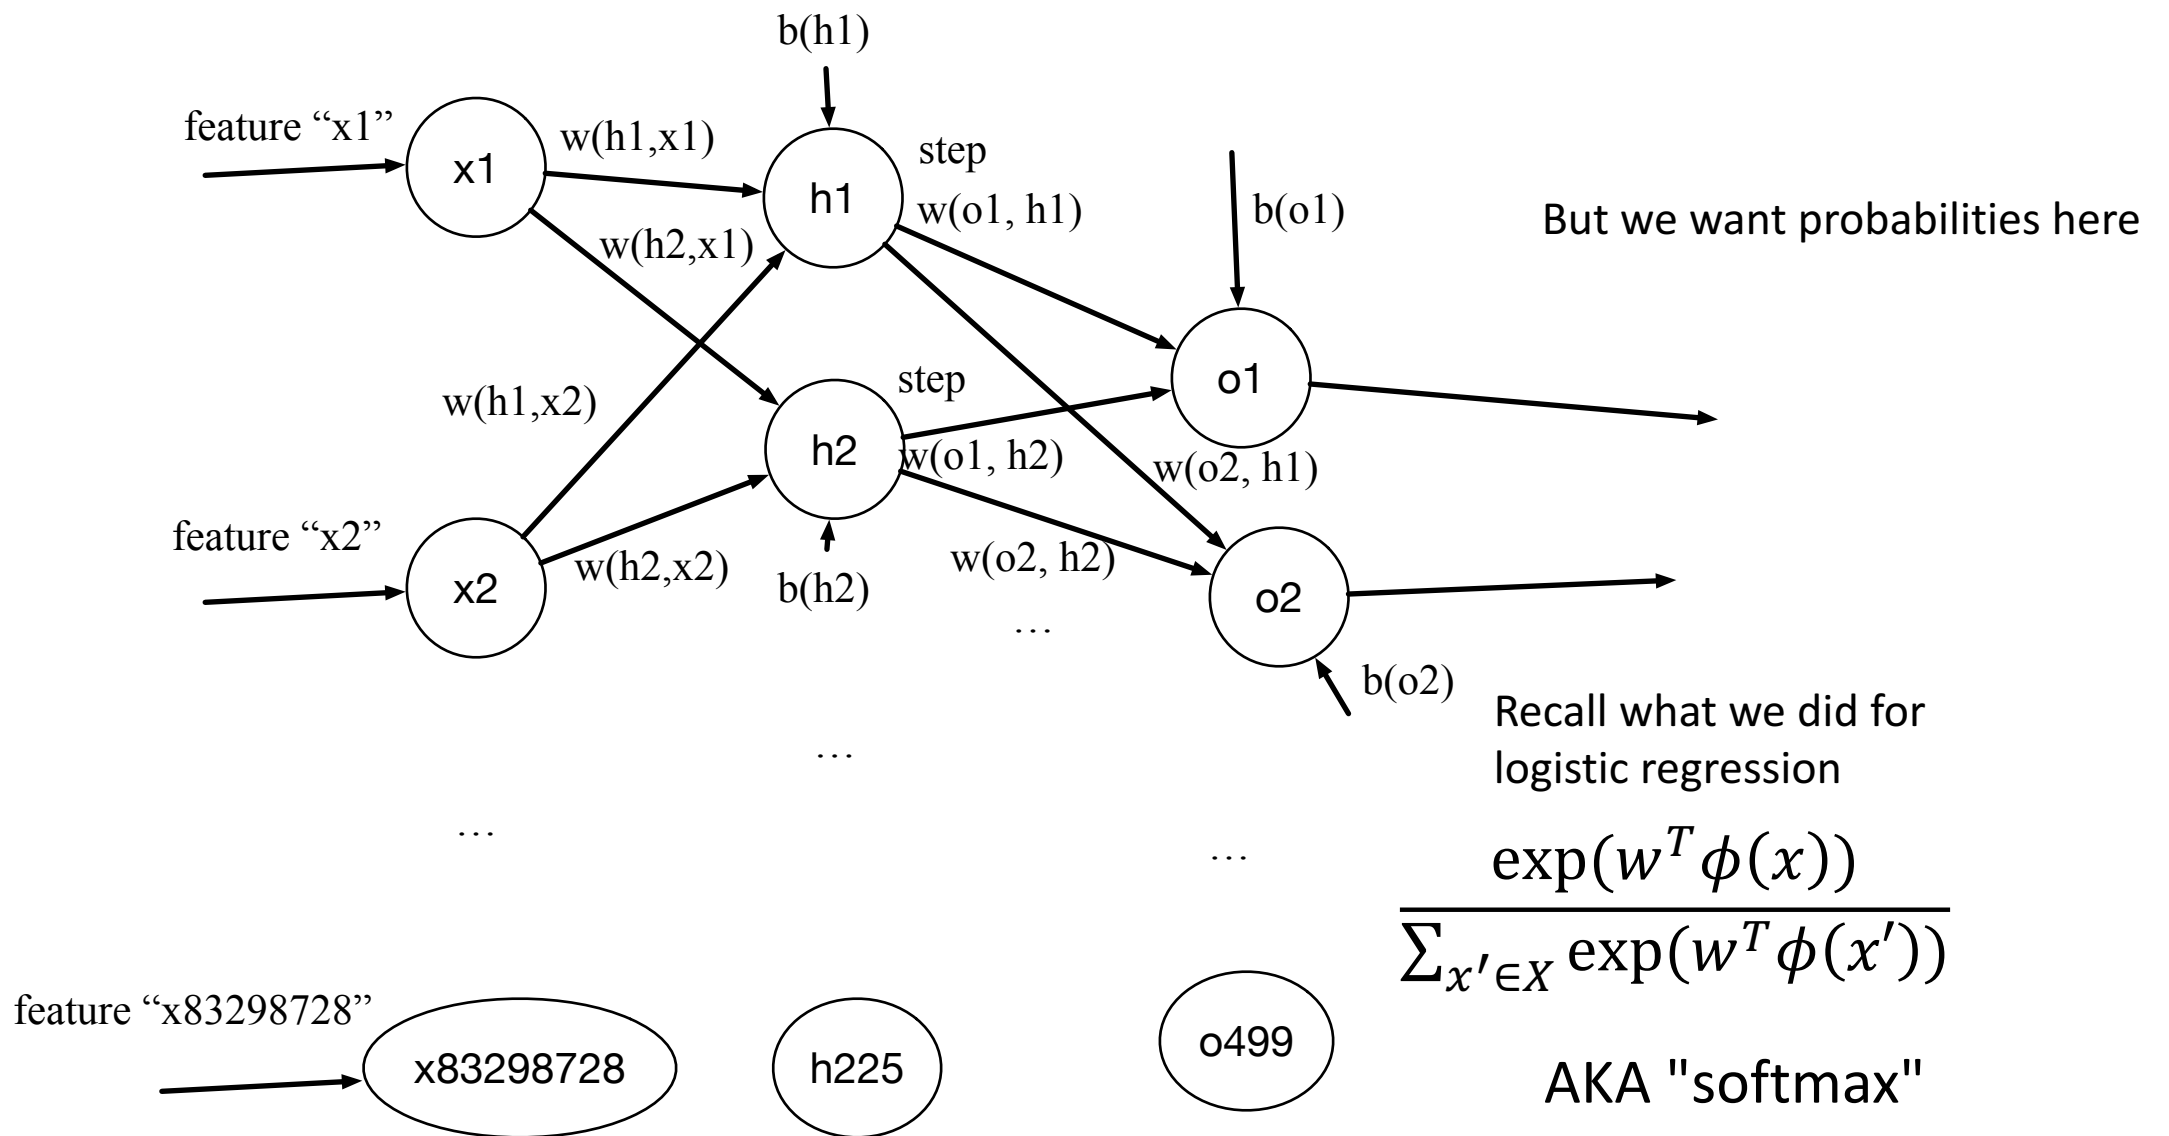
stack	symbols
[root] a b c e f g j	k l o p

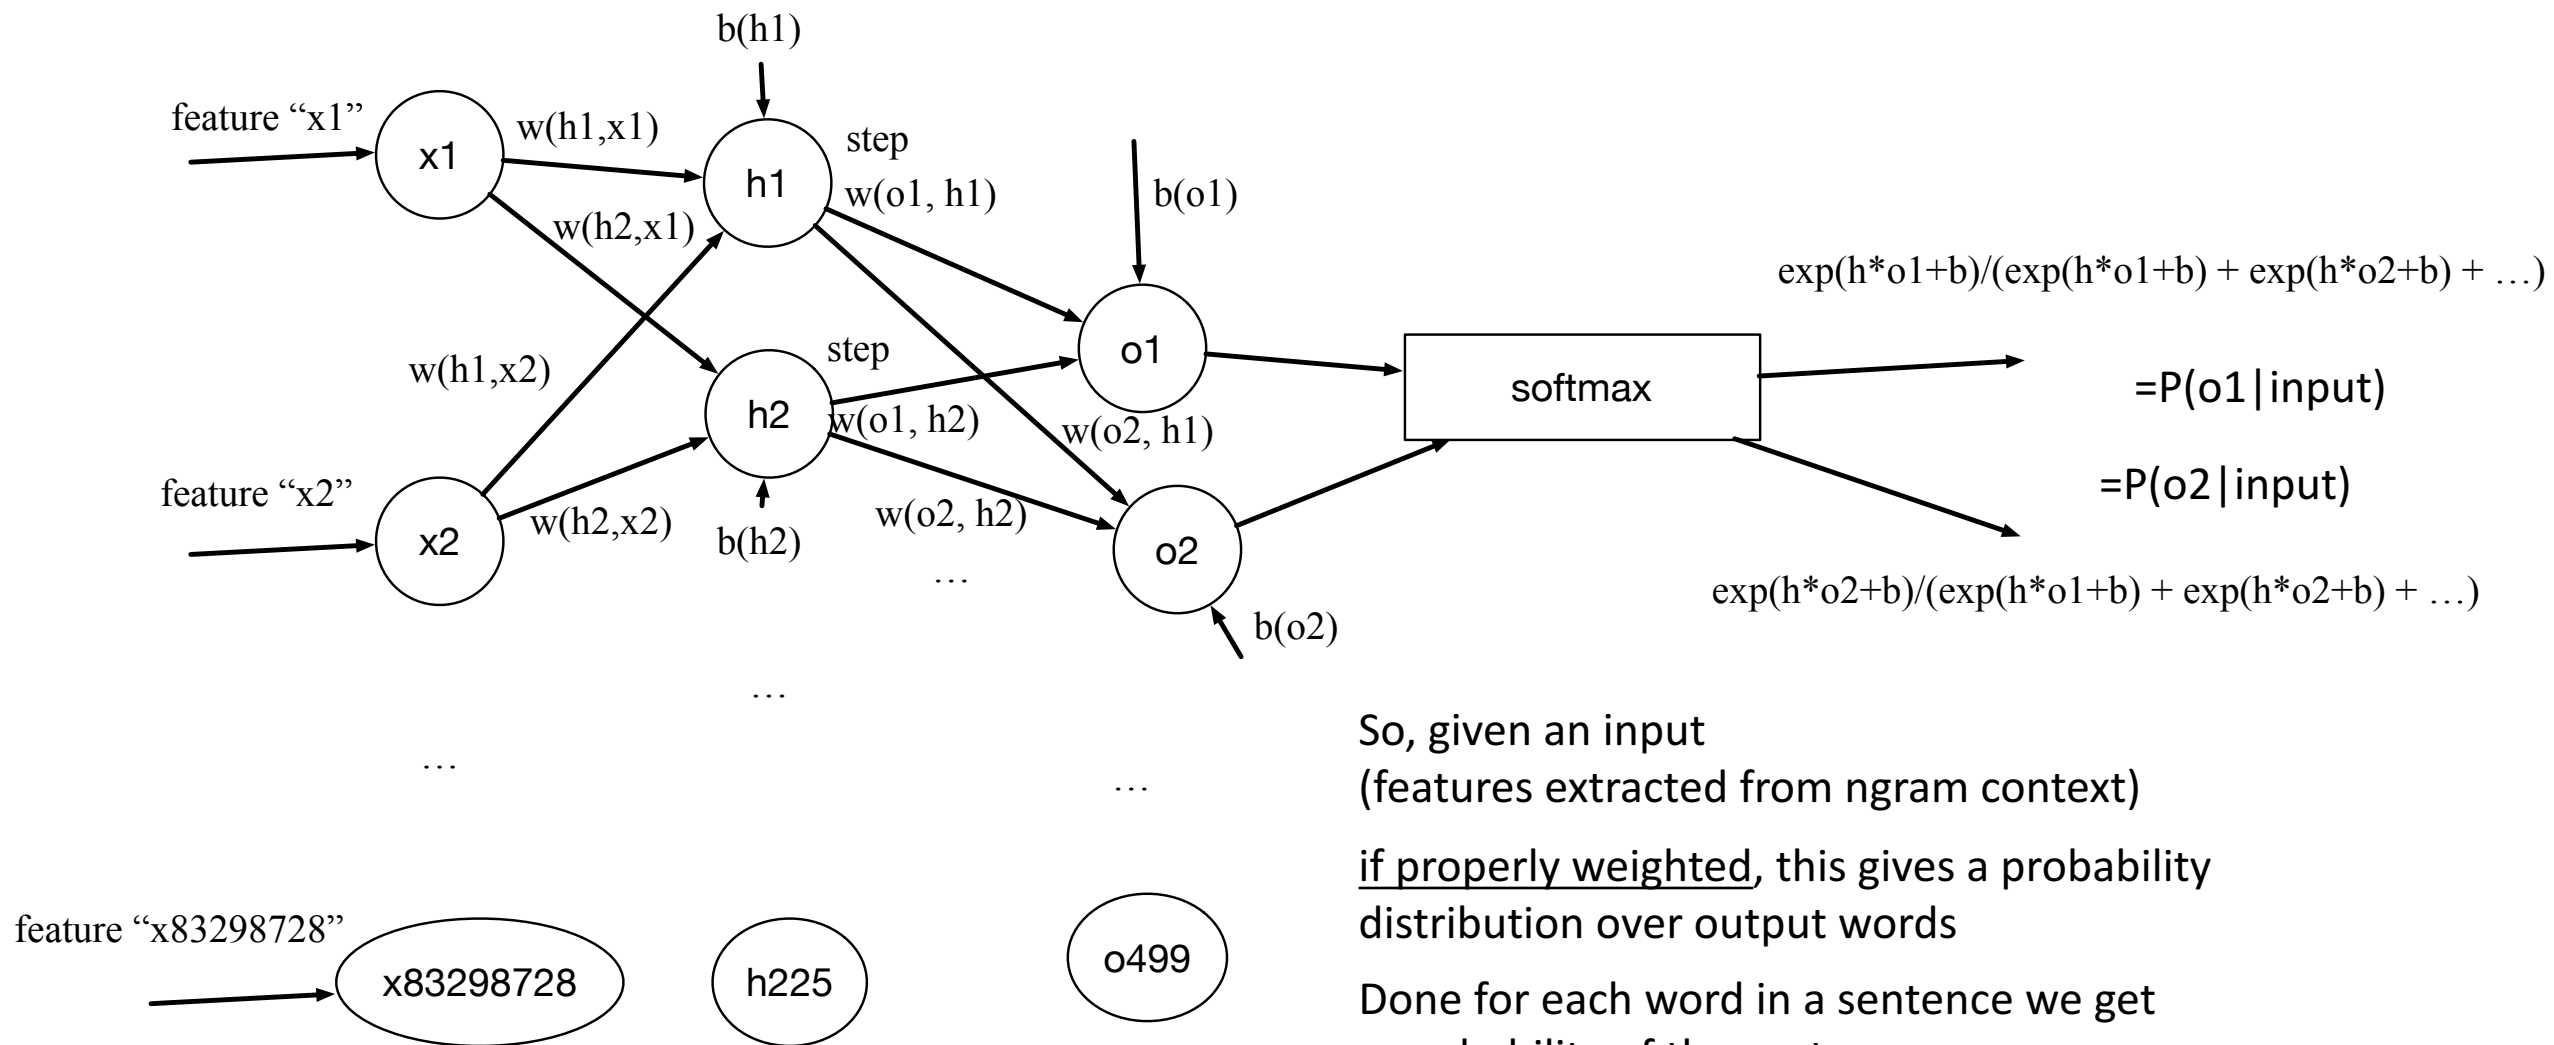
What Has This To Do With Language Models?

- Toy example:
 - x_1 (think "feature 1") can have value 1 or -1
 - x_2 can have value 1 or -1
 - outcome can be 1 or -1
 - given training data (a, b, c, d), can we use perceptron to learn o from features?
 - no. If we added x_3 (" x_1 xor x_2 ") we could. But we saw we could add layers and a nonlinear function and now use conjuncts of these features

What Has This To Do With Language Models?

- LM example:
 - x_1 can be 'wn-1=dog' (1 or 0)
 - x_2 can be 'wn-1 = the' (1 or 0)
 - ...
 - x_{24657} can be 'wn-4 is adjective'
 - we can have $x_{2947502393}$ be "wn-1=dog ^ wn-2=the" but we can also use multi-layer structure to get that for 'free'
 - for o , we'd like to know what we think about each word. So have an o for each word.





So, given an input
(features extracted from ngram context)
if properly weighted, this gives a probability
distribution over output words

Done for each word in a sentence we get
a probability of the sentence

This is a language model!

Aside: Softmax

- If you've been exposed to logistic regression and/or neural networks before you've probably heard of softmax

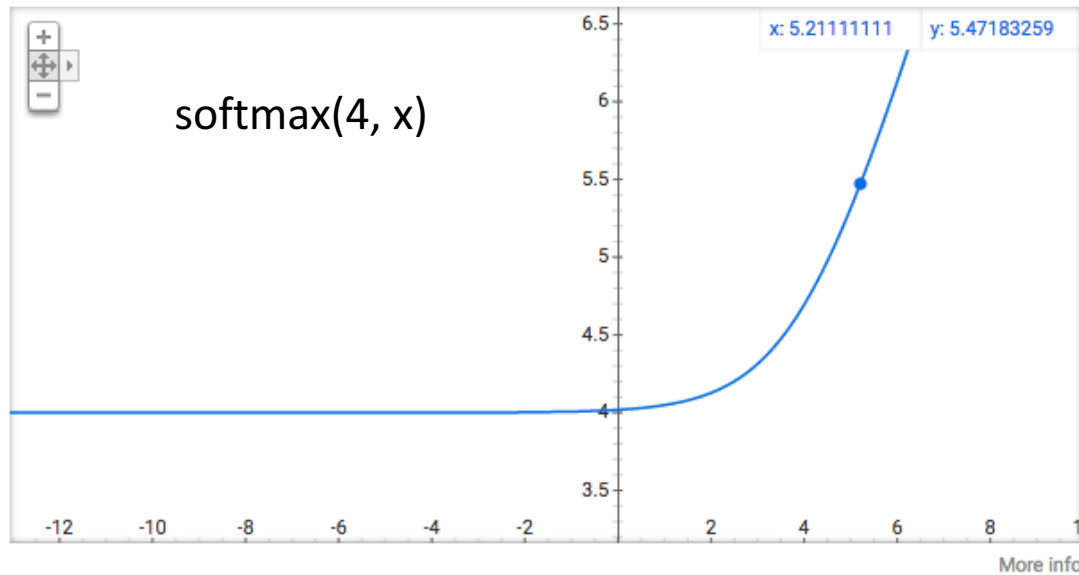
$$\frac{\exp(w^T \phi(x))}{\sum_{x' \in X} \exp(w^T \phi(x'))}$$

- I always wondered where the name comes from
- Nobody ever told me!
- Do they tell you?


Max and softmax functions

- the softmax *function* is defined as is a "soft" approximation of max
- $\text{softmax}(x, y, z) = \ln(e^x + e^y + e^z)$

Graph for $\ln(e^x + e^4)$



$\max(4, x)$



Softmax vs softmax activation

- So in general, $\text{softmax}(X) = \ln \sum_{x \in X} e^x$
- But that's not what we're usually talking about in logreg/neural network land
- We talk about $\frac{e^{x_i}}{\sum_{x \in X} e^x}$ (I simplified from the dot-product-of-features notation)
- This is useful because it squashes a collection of numbers into a probability distribution, yet preserves order
 - Remember, e^x to make everything positive, then normalize

Softmax vs softmax activation

- Consider how you might really use this, though:
 - $\frac{e^{x_i}}{\sum_{x \in X} e^x}$ is going to run into underflow issues. Best to take log
 - $\ln\left(\frac{e^{x_i}}{\sum_{x \in X} e^x}\right) = \ln e^{x_i} - \ln \sum_{x \in X} e^x$
 - $x_i - \text{softmax}(X)$
- So when we say "apply the softmax activation function" we really mean "subtract softmax from each element"

Quiz 8

stack	symbols
[root] a b c e f g	j k l o p

- What is the resulting configuration after an arc-standard Right-Arc?

stack	symbols
[root] a b c e f	g j k l o p

stack	symbols
[root] a b c e f	j k l o p

stack	symbols
[root] a b c e g	j k l o p

stack	symbols
[root] a b c f g	j l o p

stack	symbols
[root] a b c e f g j	k l o p

Quiz 9

stack	symbols
[root] a b c e f g	j k l o p

- What is the resulting configuration after an arc-eager Left-Arc?

stack	symbols
[root] a b c e f	g j k l o p

stack	symbols
[root] a b c e f	j k l o p

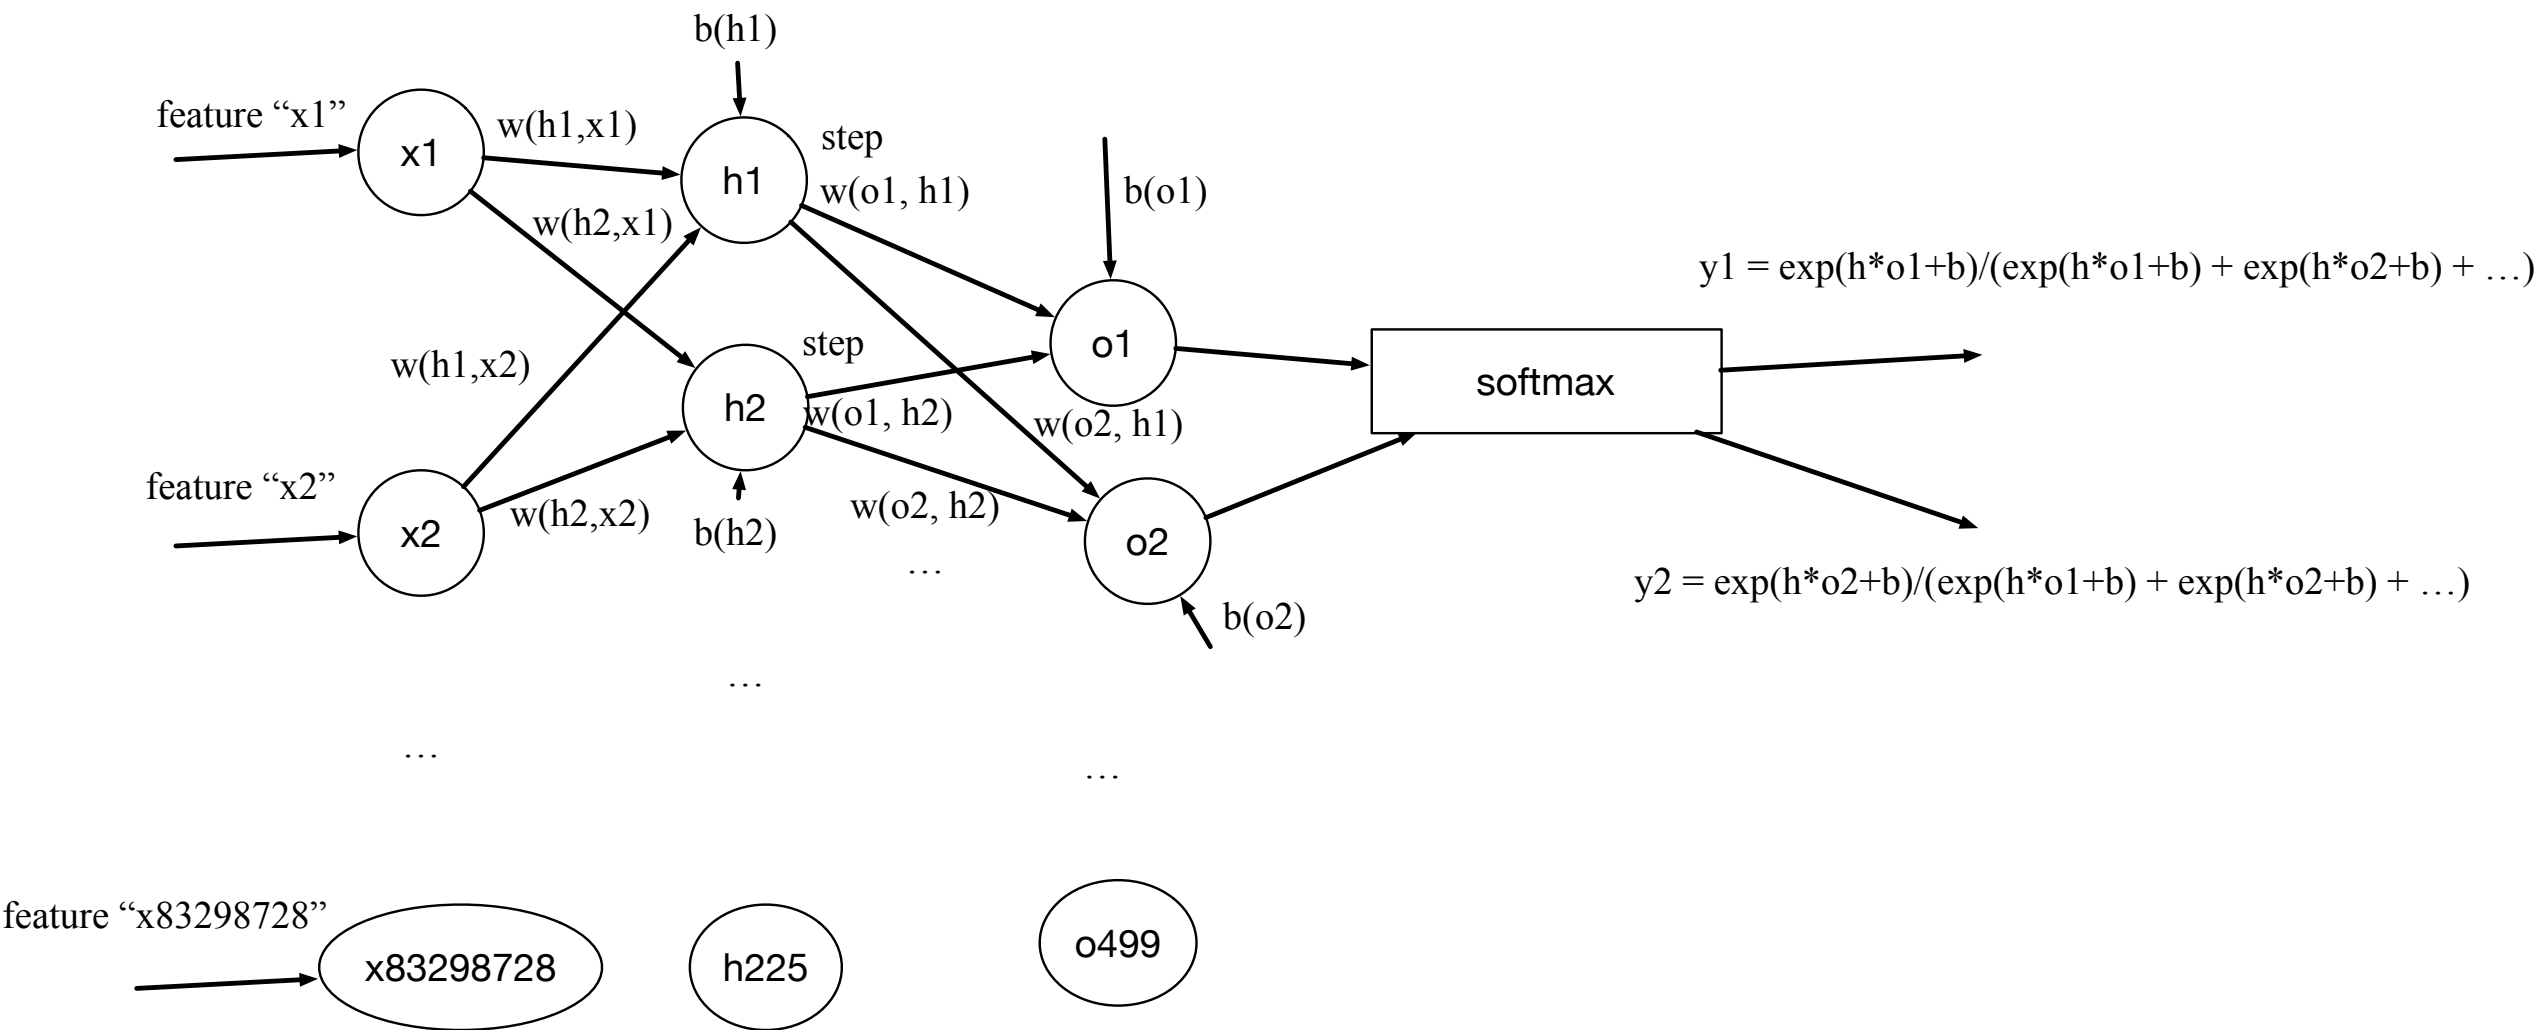
stack	symbols
[root] a b c e g	j k l o p

stack	symbols
[root] a b c f g	j l o p

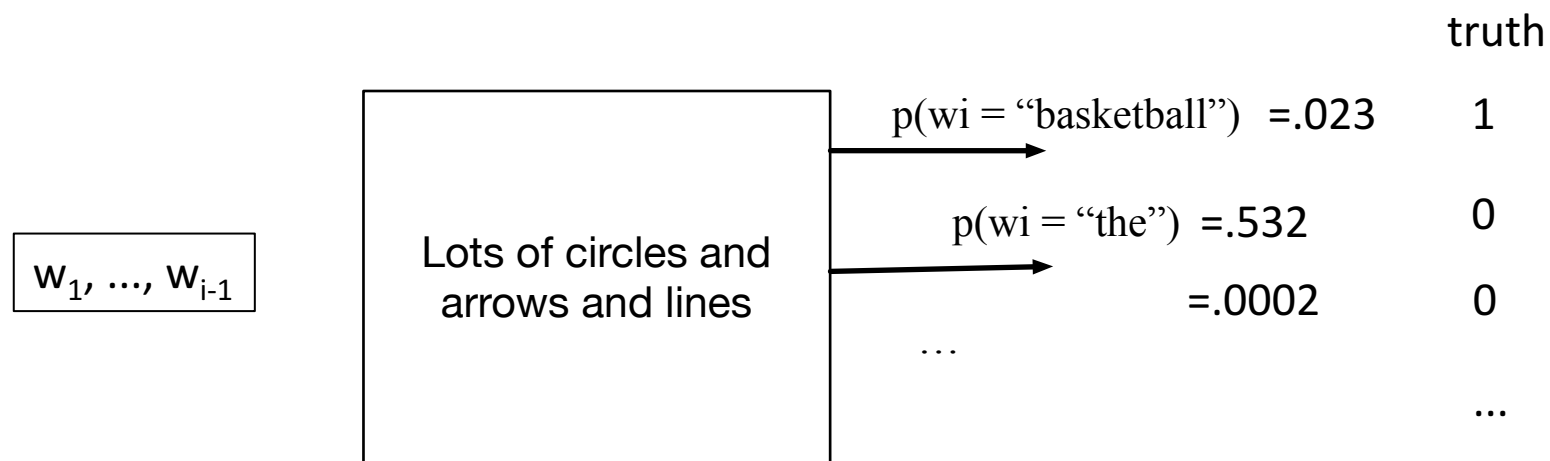
stack	symbols
[root] a b c e f g j	k l o p

How Do We Set The Weights?

- For the very simple xor case we set the 9 weights by hand
- But the general problem has lots of parameters!
- Fear not, we can use the same approach we used before:
 - Define a loss (how bad was our decision vs reality?)
 - Calculate the gradient (derivative w/r/t our parameters)
 - Adjust parameters, to move away from the gradient
 - Try again with more data, until we find something good



Training Setup



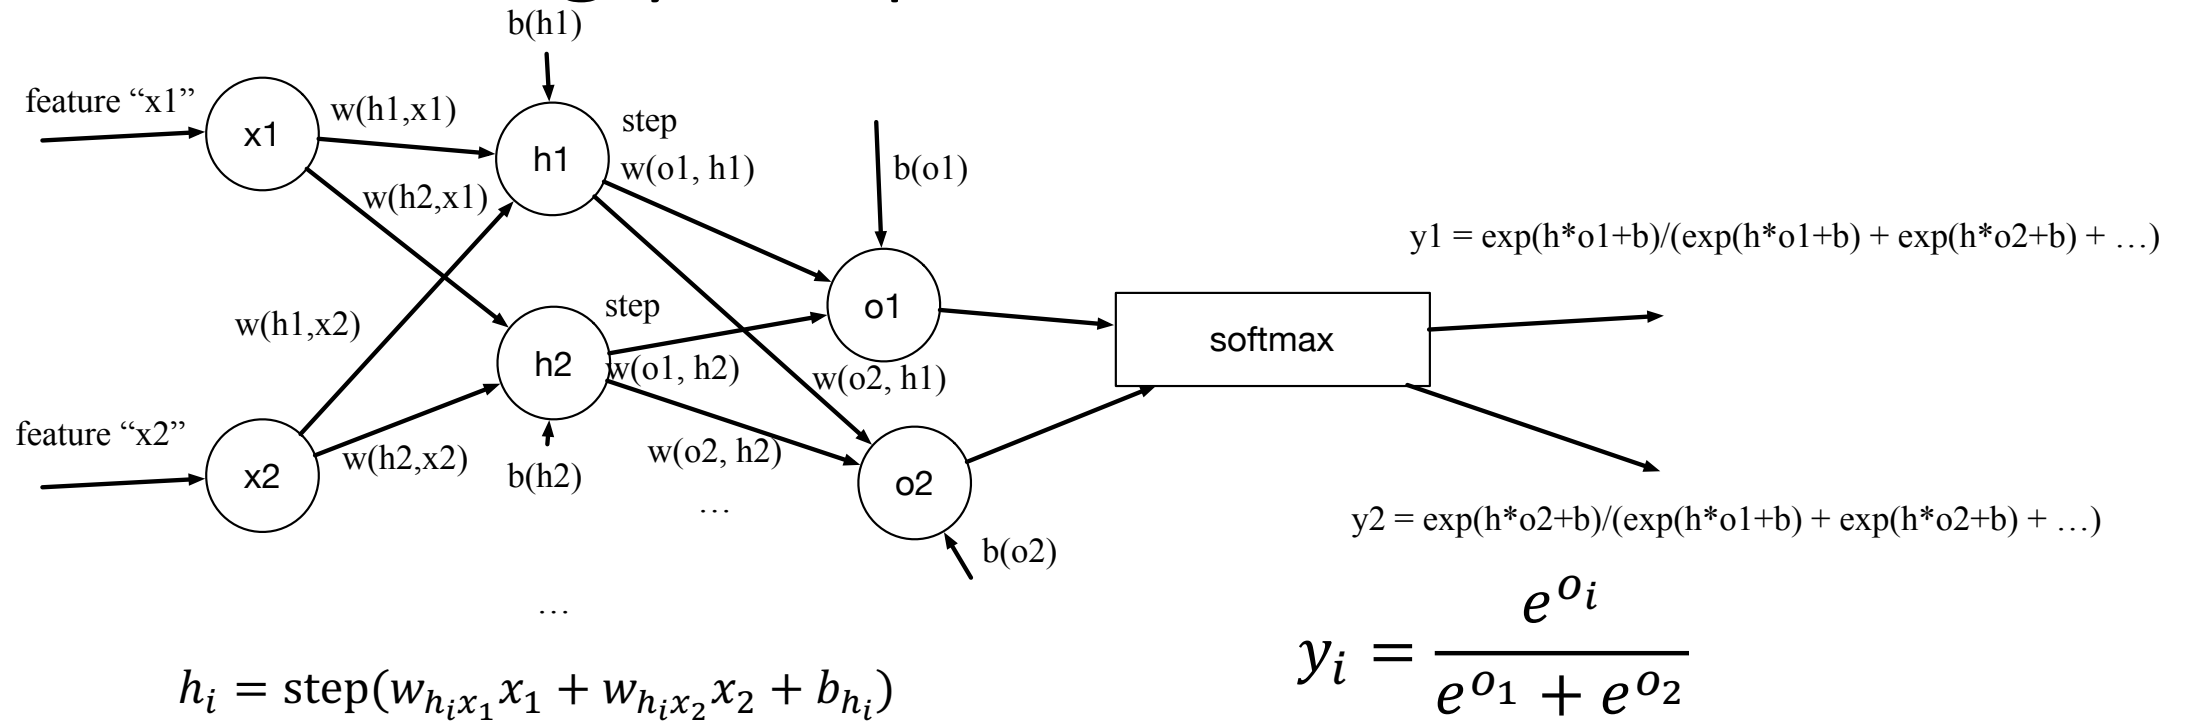
Loss

- A common one is squared error: $(y_{\text{truth}} - y_{\text{hyp}})^2$
- gradient = $2(y_{\text{truth}} - y_{\text{hyp}}) y_{\text{hyp}}'$
- y is a vector (one entry per output vocab member)
- Note: $2(y_{\text{truth}} - y_{\text{hyp}}) > 0$ for truth, < 0 else
 - = move toward the good thing, away from the bad
- Ok, so what's y_{hyp} ?

hyp	truth	
.023	1	
.532	0	
.0002	0	(scalar, given)

(result of equation;
function of parameters)

Back to the Ugly Graph



- y is pretty complicated! Need to differentiate w/r/t each parameter
- y contains step function: not differentiable at 0 and $=0$ elsewhere!
- \tanh is a nicer approximation

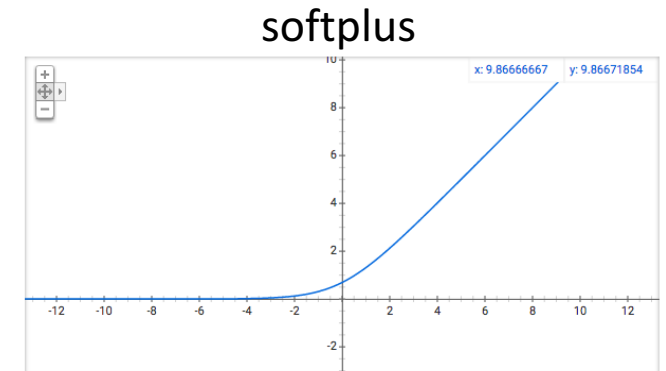

Making differentiable

$$h_i = \tanh(w_{h_i x_1} x_1 + w_{h_i x_2} x_2 + b_{h_i})$$

$$y_i = \frac{e^{o_i}}{e^{o_1} + e^{o_2}}$$

$$o_i = \tanh(w_{o_i h_1} h_1 + w_{o_i h_2} h_2 + b_{o_i})$$

- y is pretty complicated! Need to differentiate w/r/t each parameter
- y contains step function: not differentiable at 0 and =0 elsewhere!
- tanh is a nicer approximation
- Can also use ReLU (or softplus = softmax(x, 0))



Differentiating

- gradient = $2(y_{\text{truth}} - y_{\text{hyp}}) y_{\text{hyp}}'$
- y_{truth} given; y_{hyp} found by propagating data through the messy function

$$h_i = \text{step}(w_{h_1 x_1} x_1 + w_{h_2 x_2} x_2 + b_{h_i})$$

- y_{hyp}' ? lots of partials

- $dy/dw_{oh} = dy/do \cdot do/dw_{oh}$

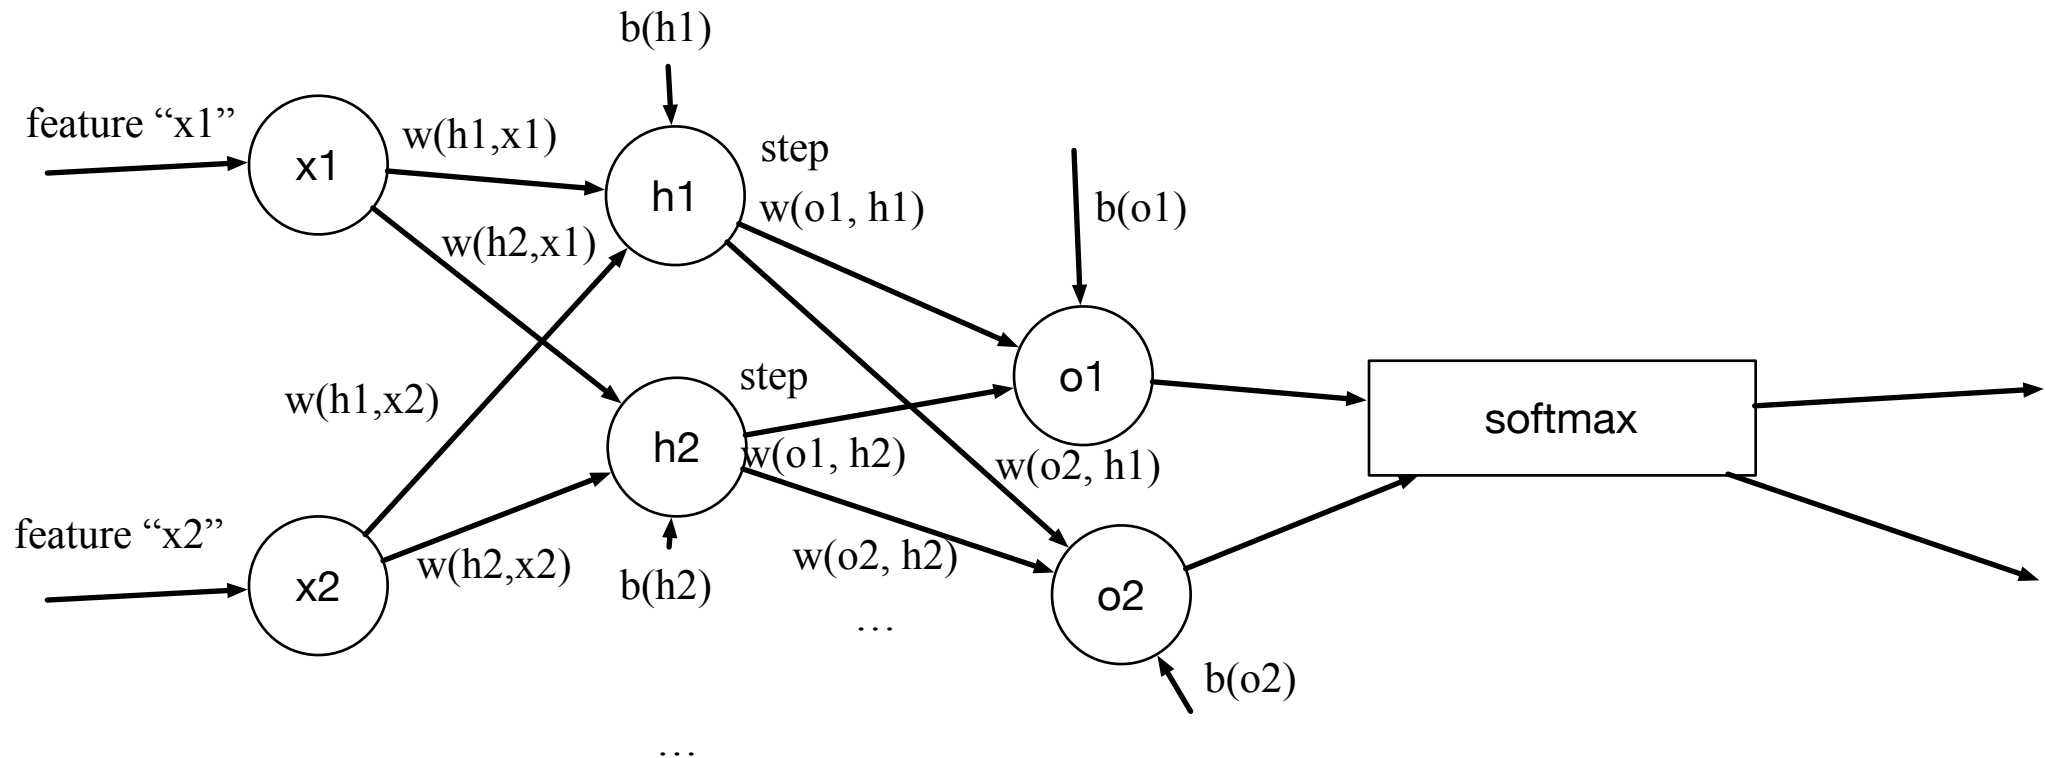
$$y_i = \frac{e^{o_i}}{e^{o_1} + e^{o_2}} \quad o_i = \text{step}(w_{o_1 h_1} h_1 + w_{o_2 h_2} h_2 + b_{o_i})$$

- $dy/db_o = dy/do \cdot do/db_o$

- $dy/dw_{hx} = dy/do \cdot do/dh \cdot dh/dw_{hx}$

- $dy/db_h = dy/do \cdot do/dh \cdot dh/db_h$

Good News: You don't really have to worry about it!



Auto-differentiation: topologically calculate values forward, derivatives backward

Partial values stored at each cell; dynamic programming makes it all efficient

Implemented in e.g. tensorflow, theano, Dynet

What Should We Connect? What Features Should We Use?

- Motivation was bigram features via structured perceptron
- So just connect the unigrams for adjacent words together?
 - i.e. all $w_1=...$ to all $w_2=...$
 - seems like a lot of careful planning
- What about similar word-class behavior?
 - Maybe all days should function similarly
 - Or all animals
- Maybe we can characterize a single word by a set of features
 - But which features?
 - Letter it starts with?
 - Part of speech?
 - Class?
- Idea: let the learning figure out how to assign features; we just choose the number of features

Fully connected

"Embeddings" shared

Embedding cell 14:
animate noun(?)

hidden cell 44:
topic is business (?)




Figure 8.13 learning all the way back to embeddings. notice that the embedding matrix E is shared among the 3 context words.

