

Lecture 5: Naive Bayes Classification

USC VSoE CSCI 544: Applied Natural Language Processing

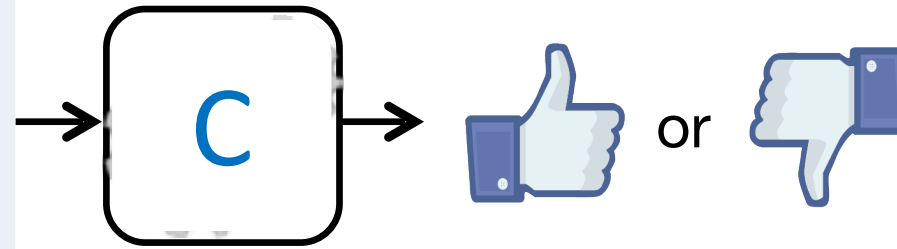
Jonathan May -- 梅約納

September 6, 2017

Sentiment Analysis

- Recall the task:

Filled with horrific dialogue, laughable characters, a laughable plot, and really no interesting stakes during this film, "Star Wars Episode I: The Phantom Menace" is not at all what I wanted from a film that is supposed to be the huge opening to the segue into the fantastic Original Trilogy. The positives include the score, the sound



- This is a **classification** task: our input is free text but our output is a fixed set of labels
- In this lecture, input/observed data is denoted x and output/prediction is y

Our Previous Rule-Based Classifier

Note: example code from lecture 2 is slightly different but functionally equivalent

```
good = { 'yay', 'love', ...}  
bad = { 'terrible', 'boo', ...}
```

← from our pos/neg word list!

```
score = 0  
for w in x:  
    if w in good:  
        score +=1  
    elif w in bad:  
        score -=1  
if score >= 0:  
    return 'pos'  
return 'neg'
```

x →

C

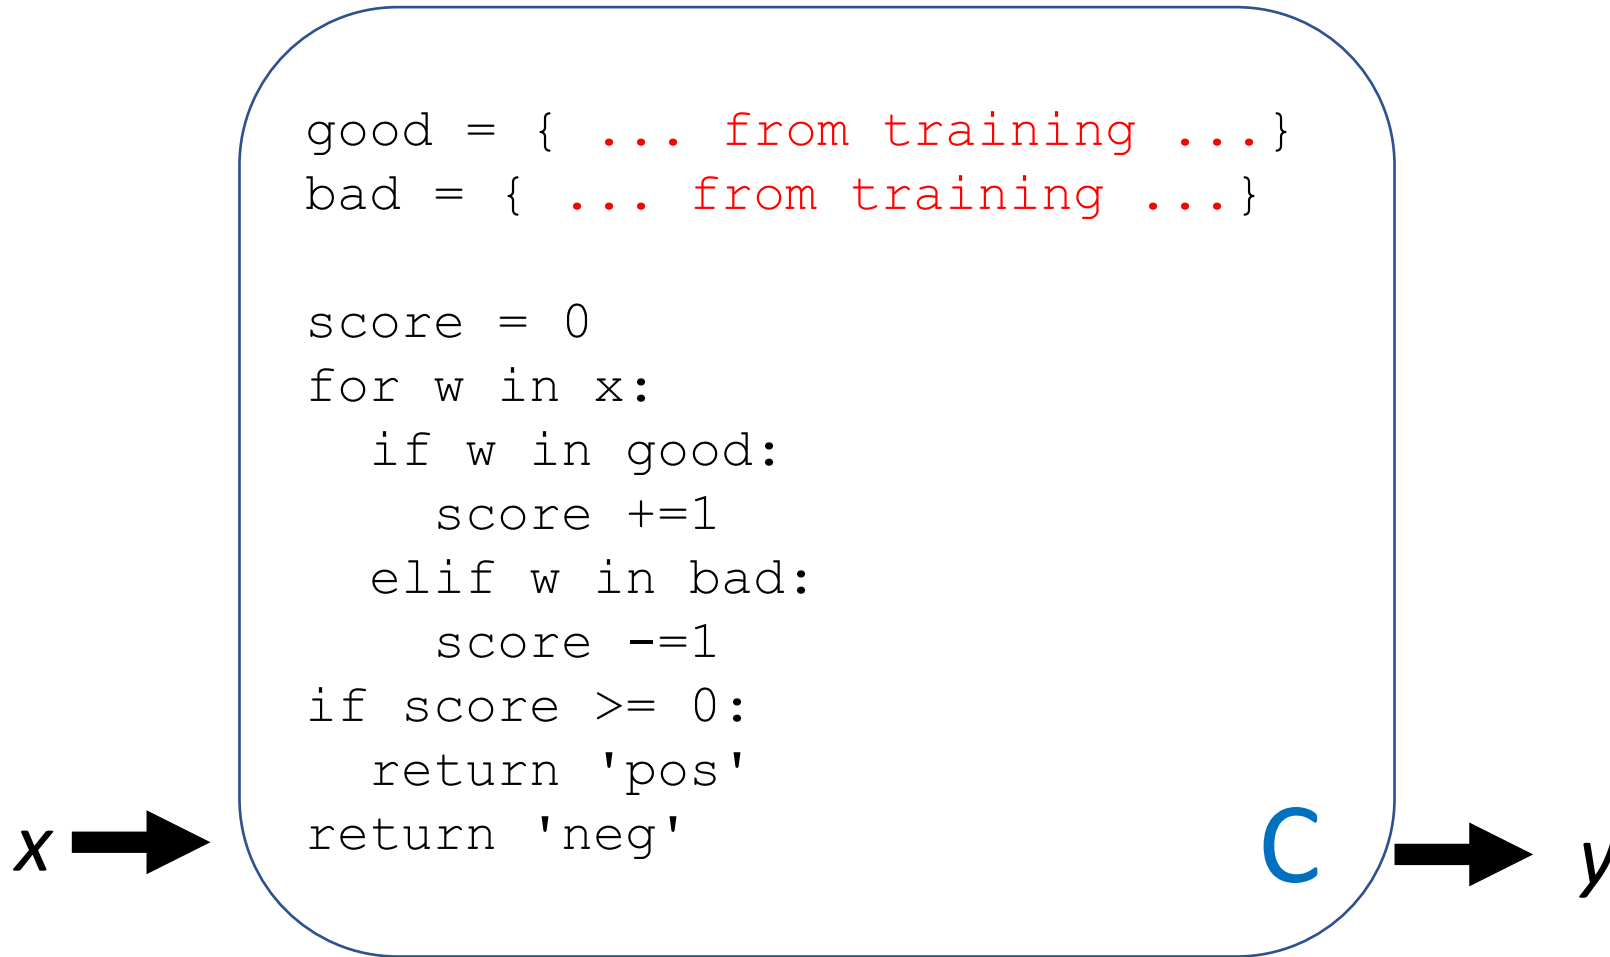
→ y

Supervised Classification

- Rather than top-down features, let's use data-driven features
 - Our intuitions about word sentiment aren't perfect
- **Supervised** (aka Inductive) = learn from **labeled** examples. Recipe:
 - **training** corpus of (x, y) (review, label) pairs
 - learning algorithm
- Other kinds of learning (not covered here):
 - **Unsupervised**: Data is provided but no labels are given
 - **Semi-Supervised**: Data is provided but only some of it is labeled
 - **Reinforcement**: No clear labels, but feedback (in the form of a reward) is accessible

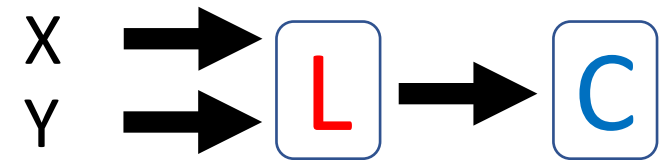

Supervised

Our ~~Previous Rule Based~~ Classifier



Notation

- Training examples: $X = (x_1, x_2, \dots, x_N)$
- Labels of training examples: $Y = (y_1, y_1, \dots, y_N)$
- A classifier **C** maps x_i to y_i
- A learner **L** infers **C** from (X, Y)



Counting-based Learner

X →
Y →

```
from collections import Counter
scores = Counter()
for x, y in zip(X, Y):
    for w in x:
        if y == 'pos':
            scores[w]+=1
        elif y == 'neg':
            scores[w]-=1
good, bad = set(), set()
for w, score in scores.items():
    if score>=0: good.add(w)
    else: bad.add(w)
return good, bad
```

L



C

Probability Review Questions

- 1) If P is a probability function, which of the following is equal to $P(x \mid y, z)$?
 - A) $P(x) / P(y, z)$
 - B) $P(y)P(z) / P(x, y, z)$
 - C) $P(x, y, z) / P(y, z)$
 - D) $P(x)P(x \mid y)P(x \mid z)$

Probability Review Questions

- 2) Which is/are guaranteed to be true?
 - A) $\forall y \forall z, \sum_x p(x | y, z) = 1$
 - B) $\forall x, \sum_y \sum_z p(x | y, z) = 1$
 - C) $\sum_x p(x) = 1$
 - D) $\forall y \forall z, \sum_x p(x) p(y|x) p(z|x, y) = 1$

A Probabilistic Classifier

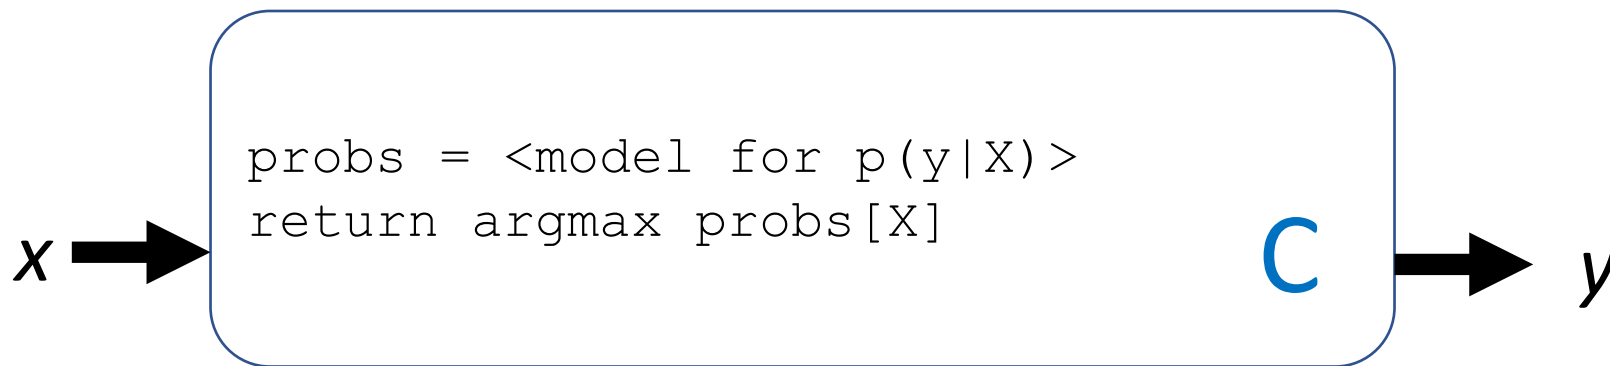
Experiment: "people wrote reviews of movies"

Outcomes (Ω): all possible reviews

Events: $Y=\text{pos}$: all positive reviews. $Y=\text{neg}$ = all negative reviews

$X=34925$: "when review 34925 was written"

Review #34925: "Filled with horrific dialogue, laughable characters, a laughable plot, and really no interesting stakes during this film, "Star Wars Episode I: The Phantom Menace" is not at all what I wanted from a film that is supposed to be the huge opening to the segue into the fantastic Original Trilogy. The positives include the score, the sound..."



MLE is not going to work (review 34925 shouldn't be in training)!

Aside: Library of Babel (Borges, 1941)

- Contains all 1.3m-word texts
- A lot of it looks like junk
- But all useful texts are in here too!
- What's it like to be a librarian?
- Look up your favorite piece of text at <https://libraryofbabel.info>



A probabilistic model that generalizes

- Instead of estimating $p(Y \mid \text{Filled with horrific...})$ directly, let's make two **modeling assumptions**:
 1. The **Bag of Words (BoW) assumption**: Assume the order of the words in the document doesn't matter:
 $p(Y \mid \text{Filled with horrific ...}) = P(Y \mid \text{with, horrific, Filled, ...})$

a sequence

independent word events

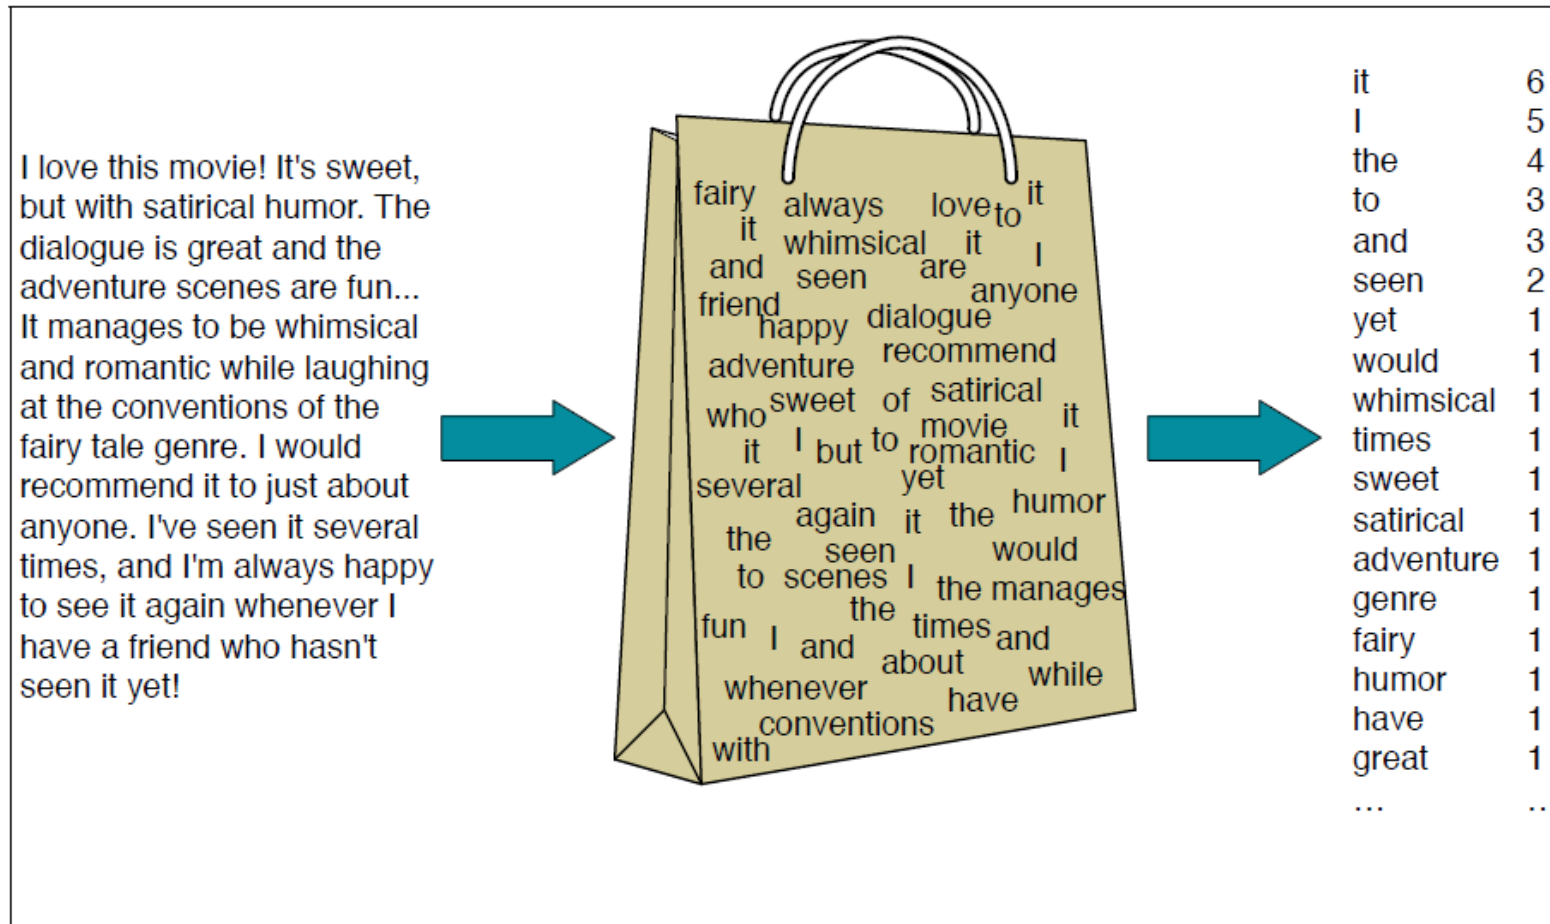


Figure 7.1 Intuition of the multinomial naive Bayes classifier applied to a movie review. The position of the words is ignored (the *bag of words* assumption) and we make use of the frequency of each word.

Figure from J&M 3rd ed. draft, sec 7.1

A probabilistic model that generalizes

- Instead of estimating $p(Y \mid \text{Filled with horrific...})$ directly, let's make two **modeling assumptions**:
 1. The **Bag of Words (BoW) assumption**: Assume the order of the words in the document doesn't matter:
 $p(Y \mid \text{Filled with horrific ...}) = P(Y \mid \text{with, horrific, Filled, ...})$

a sequence

independent word events

So called because a **bag** or **multiset** is a data structure that stores counts of elements, but not their order

A probabilistic model that generalizes

- The BoW assumption isn't enough unless you happen to have seen one of 34925's anagrams in your training data (e.g. 4088293). Hence:
 2. The **naive Bayes assumption**: Words are independent conditioned on their class:
$$P(\text{Filled with horrific...} \mid Y) = P(\text{Filled} \mid Y) \times P(\text{with} \mid Y) \times P(\text{horrific} \mid Y) \dots$$
- Wait, but we were estimating $P(Y \mid \text{with, horrific, Filled, ...})$. What should we do?
- Bayes Rule!

Quiz

- Which of the following is/are equal to $P(A|B, C, D)$?
 - 1) $P(B, C, D | A) P(A)$
 - 2) $P(B|C, D) P(B|A, C) P(B | A, D)$
 - 3) $P(D, C, B | A) P(A) / P(D, C, B)$
 - 4) $P(A | D, C, B)$
- Note that while $P(A|B, C, D) \neq P(B, C, D | A) P(A)$,
 $\operatorname{argmax}_A P(A|B, C, D) = \operatorname{argmax}_A P(B, C, D | A) P(A)$, and that's what we care about.

A probabilistic model that generalizes

- Put the assumptions/rule together:
 - $p(Y | \text{Filled with horrific ...}) = P(Y | \text{with, horrific, Filled, ...})$ (BoW assmptn)
 - $\propto P(Y) \times P(\text{with, horrific, Filled, ...} | Y)$ (Bayes' rule)
 - $= P(Y) \times P(\text{with} | Y) \times P(\text{horrific} | Y) \times P(\text{Filled} | Y) \dots$
(Naive Bayes assmptn)

Is this a good model?

- "all models are wrong, but some are useful" – George Box, statistician
- What's wrong with the BoW assumption?
- What's wrong with the Naive Bayes assumption?
- But does it work?
 - Yes, for many tasks
 - And that's kind of all that matters

Naive Bayes Classifier

```
from numpy import argmax
wprobs = { from training }
cprobs = { from training }
totals = []
for c in classes:
    total = class_probs[c]
    for w in x:
        total *= wprobs[w][c]
    totals.append(total)
return argmax(totals)
```

x →

C

→ y

Naive Bayes Learner

X →
Y →

```
from collections import Counter
cscores = Counter()
wscores = []
for c in classes:
    wscores.append(Counter())
for x, y in zip(X, Y):
    cscores[y] += 1
    for w in x.split():
        wscores[y][w] += 1
cprobs, wprobs = [], []
for c in classes:
    cprobs = cscores[c] / len(Y)
    wprob = {}
    for w, score in wscores[c].items():
        wprobs[w] = score / cscores[c]
    wprobs.append(wprob)
return cprobs, wprobs
```

L



C

Parameters

- Every probability or other value that is **learned** and used by the classifier is called a **parameter** (e.g. everything the learner sent the classifier)
- Naive Bayes has two kinds of parameters:
 - Class prior distribution $\mathbf{P(Y)}$ = belief in the class
 - Likelihood distribution $\mathbf{P(W|Y)}$ = likelihood of observing a word in a class
- If there are K classes and V words, how many parameters are in a Naive Bayes classifier?

Practicalities: Smoothing

- Recall:

```
for w in x:  
    total *= wprobs[w][c]
```

- What if you see a new word, or word unassociated with that class, at test time?
- Whole probability will be 0!

Laplace (add-1) smoothing

- Assume we've seen a special symbol called OOV (out of vocabulary) once per class. And assume we've seen all possibilities once more.
- Before: $p(\text{wonderful} \mid \text{pos}) = \text{count}(\text{wonderful}, \text{pos}) / \text{count}(*, \text{pos})$
 $p(\text{terrible} \mid \text{pos}) = \text{count}(\text{terrible}, \text{pos}) = 0 / \text{count}(*, \text{pos}) = 0!$

vocabulary	pos	neg
wonderful	398	17
terrible	0	228
...
TOTAL (5000 types)	147808	167585

Laplace (add-1) smoothing

- Assume we've seen a special symbol called OOV (out of vocabulary) once per class. And assume we've seen all possibilities once more.
- After: $p(\text{wonderful} \mid \text{pos}) = \frac{[\text{count}(\text{wonderful}, \text{pos}) + 1]}{[\text{count}(*, \text{pos}) + |V| + 1]}$
 $p(\text{terrible} \mid \text{pos}) = p(\text{OOV} \mid \text{pos}) = \frac{1}{[\text{count}(*, \text{pos}) + |V| + 1]}$

vocabulary	pos	neg
wonderful	398	17
terrible	0	228
...
TOTAL (5000 types)	147808	167585

vocabulary	pos	neg
wonderful	399	18
terrible	1	229
...	...+1	...+1
OOV	1	1
TOTAL (5000 types)	147808+5001	167585+5001


Practicalities: Underflow

- Recall:
for w in x:
 total *= wprobs[w][c]
- x may be long! wprobs[w][c] may be small!
- But remember log math!
 - $\exp(\log(x)) = x$
 - $\log(xy) = \log(x) + \log(y)$
 - $\forall x, y > 0, x > y \Leftrightarrow \log(x) > \log(y)$

```
a=1
for i in range(100):
    a*=.0001
    if i % 10 == 0:
        print(i, a)
```

```
0 0.0001
10 1.0000000000000003e-44
20 1.0000000000000007e-84
30 1.00000001e-124
40 1.00000015e-164
50 1.00000021e-204
60 1.00000029e-244
70 1.00000036e-284
```

Graph for $\log(x)$



Avoiding Underflow with Logs

- Now:

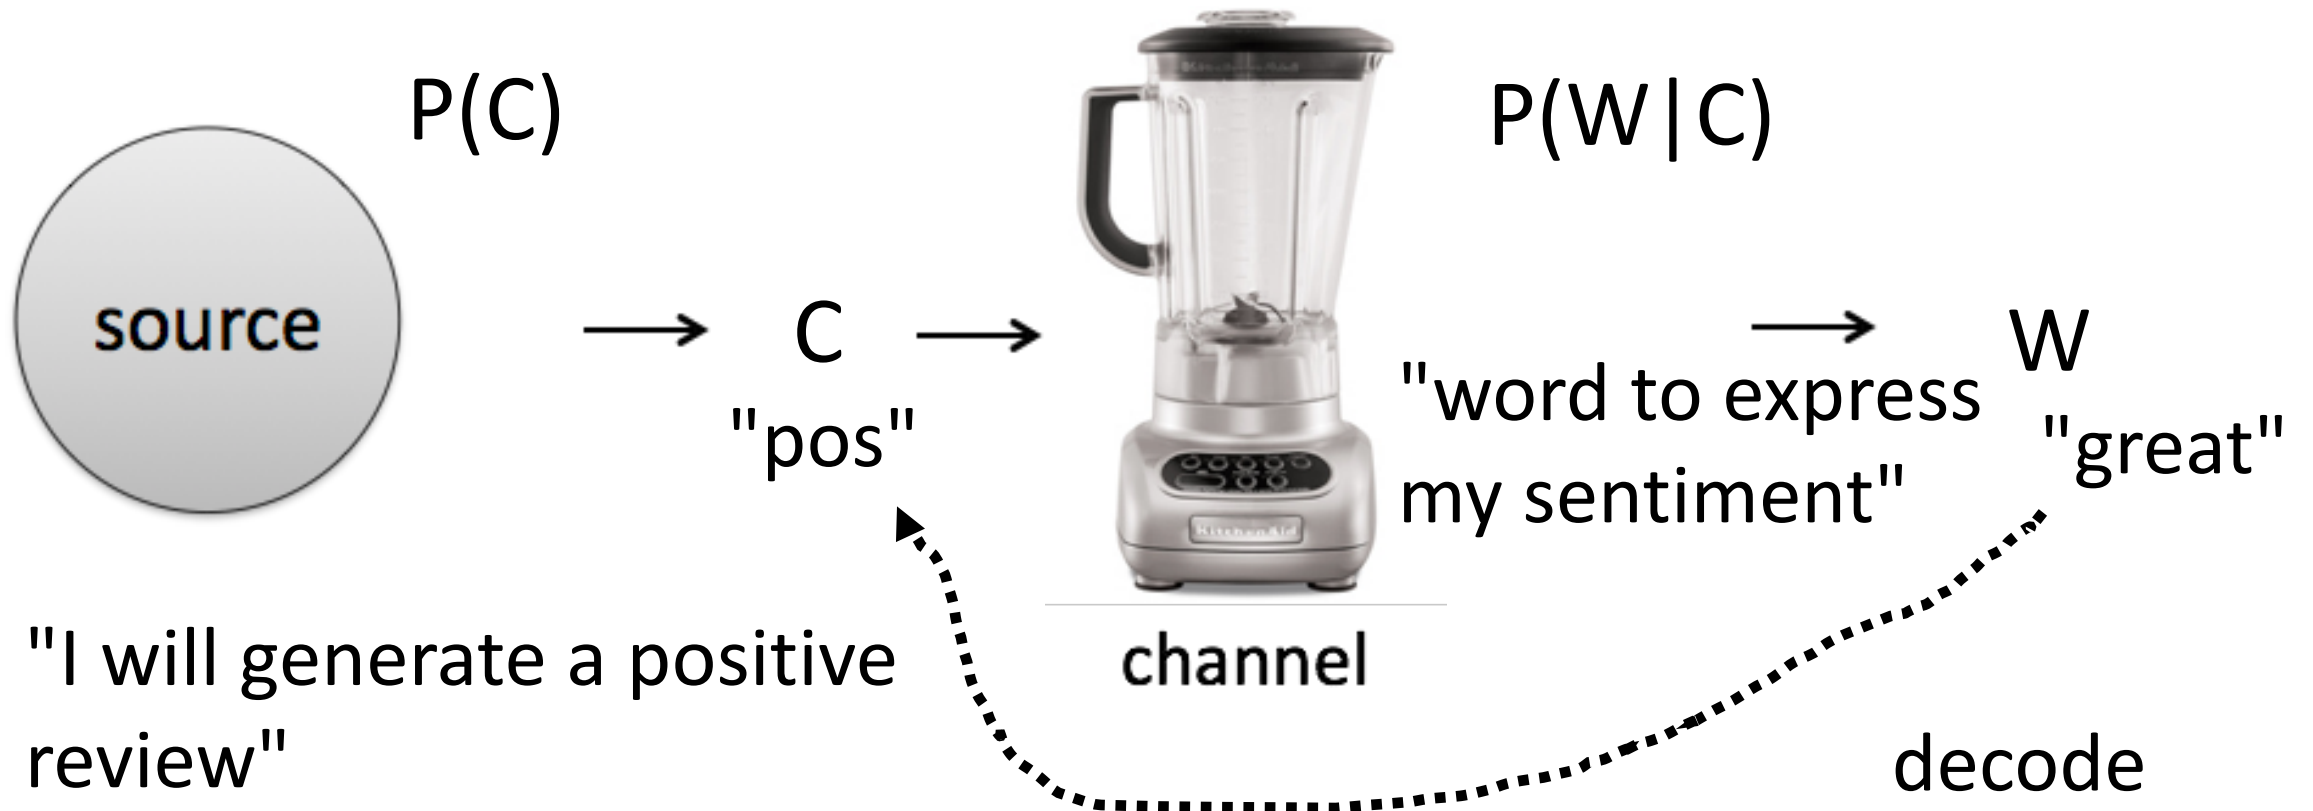
```
for w in x:  
    total += log(wprobs[w][c])
```

```
from math import log  
a = log(1)  
for i in range(100):  
    a += log(.0001)  
    if i % 10 == 0:  
        print(i, a)
```

```
0 -9.210340371976182  
10 -101.31374409173802  
20 -193.41714781149977  
30 -285.5205515312616  
40 -377.62395525102363  
50 -469.72735897078564  
60 -561.8307626905473  
70 -653.9341664103088  
80 -746.0375701300702  
90 -838.1409738498317
```

Noisy Channel Model

- Reminder: Using Bayes' Rule interpretation of the world



Conclusions

- We have seen how **training data** and **supervised learning** can produce a better classifier
 - **Classifier** takes an *input* (such as a text document) and predicts an *output* (such as a class label)
 - **Learner** takes *training data* and produces (statistics necessary for) the classifier

Conclusions

- Because most pieces of text are unique, it's not very practical to assume the one being classified is in the training data
 - though it is in the library of Babel!
 - We need to make **modeling assumptions** that help the learner to **generalize** to unseen inputs
- The **Naive Bayes** model and **Bag of Words** assumption are a simple, fast probabilistic approach to text classification